

CH2. Mathematical Building Blocks of Deep-learning

Il-Youp Kwak
Chung-Ang University



Mathematical Building Blocs of Deep-learning

A first look on neural network basic example & construction Tensorflow Keras

Data representations for neural network

The gear of neural networks: tensor operations

The engine of neural networks: gradient-based optimization

Looking back at our first example



A first look on neural network



Figure 2.1 MNIST sample digits

Mnist data analysis example

hand written digit data set

Hand written digit (28,28 pixcels) for 10 categories (0-9)

60,000 training set and 10,000 test set



* There is keras module in Tensorflow. (tf)

Mnist data

already in tensorflow module

- 케라스는 텐서플로와 별개의 라이브러리, 텐서플로에서도 케라스 사용

- In keras module, there is some widely popular used dataset.
(MNIST is one of that)

Listing 2.1 Loading the MNIST dataset in Keras

```
[from tensorflow.keras.datasets import mnist] → { calling } library, we load mnist data.
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

Try these codes from <https://colab.research.google.com/github/>

Examples GitHub

Enter a GitHub URL or search by organization or user Include private repos

fchollet

Repository: [fchollet/deep-learning-with-python-notebooks](#) Branch: [master](#)

Path

[chapter02_mathematical-building-blocks.ipynb](#) [chapter03_introduction-to-keras-and-tf.ipynb](#)



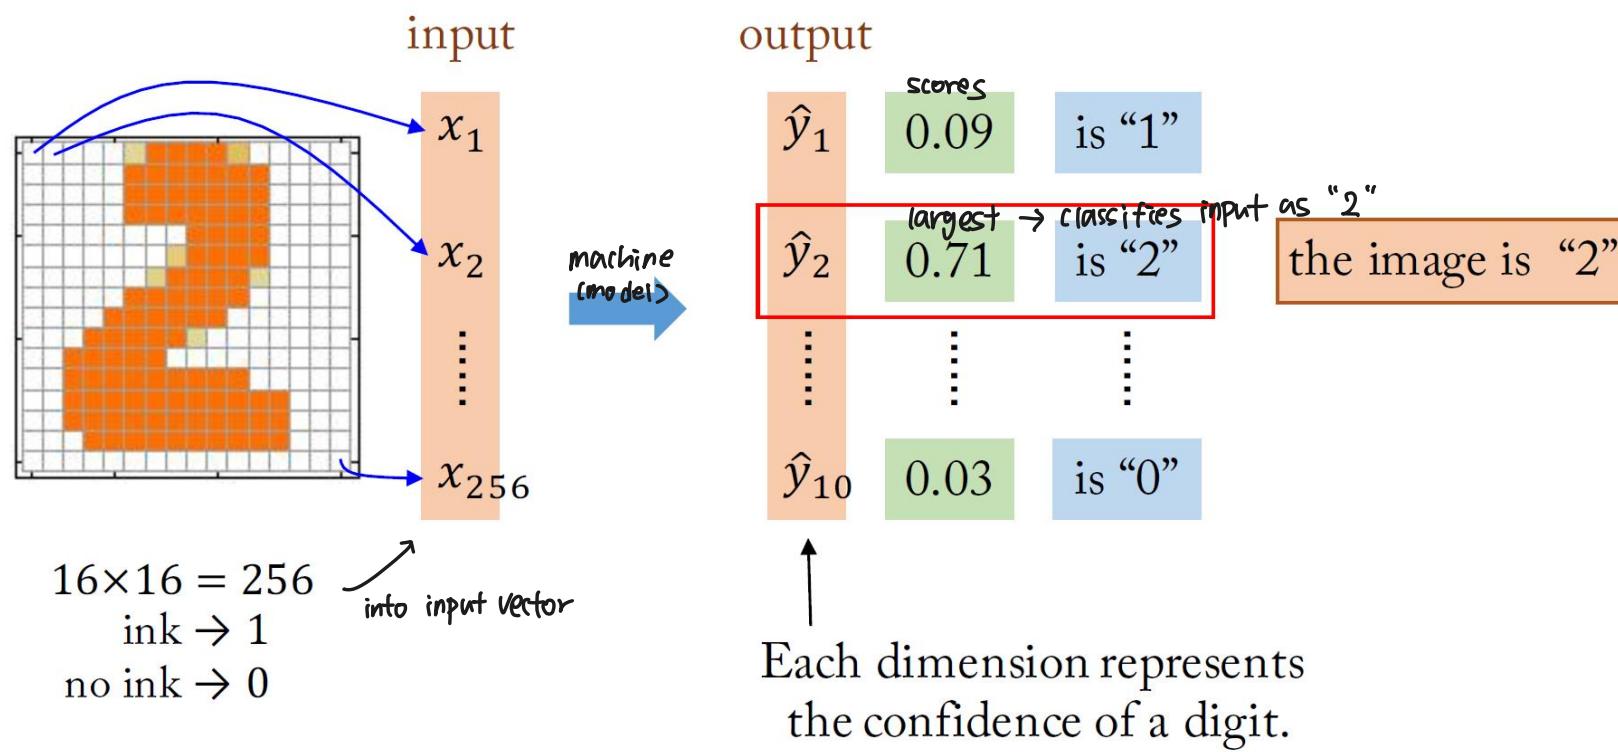
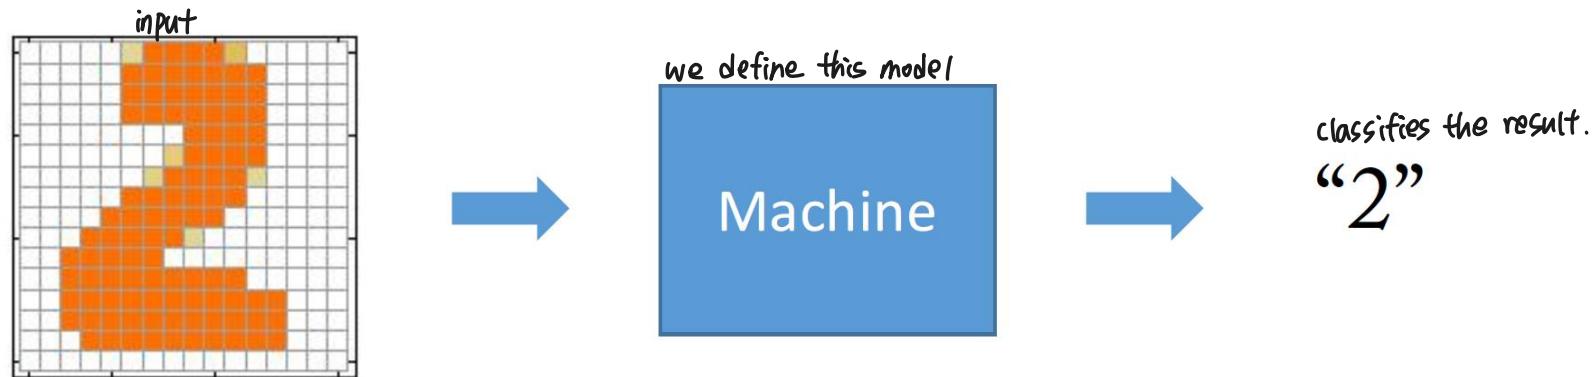
Welcome To Colaboratory

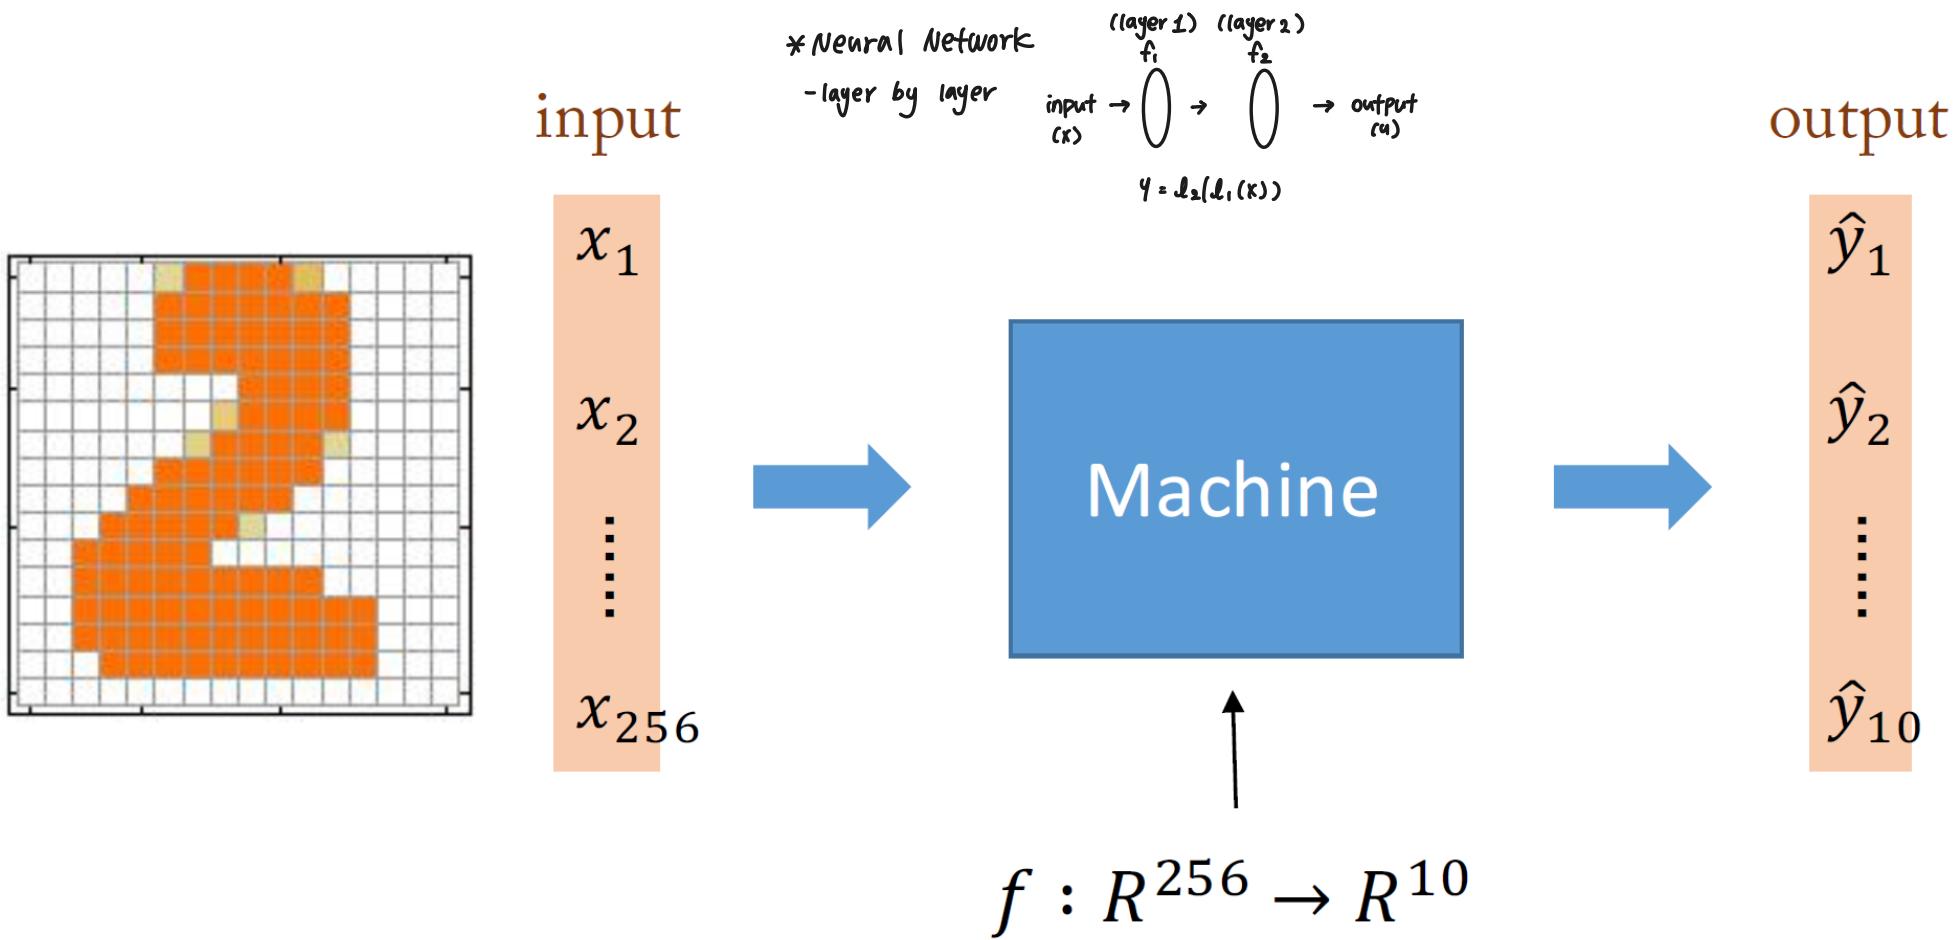
File Edit View Insert Runtime Tools Help

Recent			
Google Drive			
GitHub			
Upload			
<input type="text" value="Filter notebooks"/> ≡			
Title	Last opened	First opened	⋮
 Welcome To Colaboratory	12:29 PM	Sep 23, 2019	
 3.5-classifying-movie-reviews.ipynb	March 4	Mar 14, 2021	 
 2.1-a-first-look-at-a-neural-network.ipynb	March 4	Feb 27, 2021	 
 chapter02_mathematical-building-blocks.ipynb	March 3	January 24	 
 neurokit2.ipynb	February 28	February 28	 

Whether you're a **student**, a **data scientist** or an **AI researcher**, Colab can make your work easier. Watch [Introduction to Colab](#) to learn more, or just get started below!

Ch2. Mathematical building blocks of deep-learning / 2.1 A first look on neural network

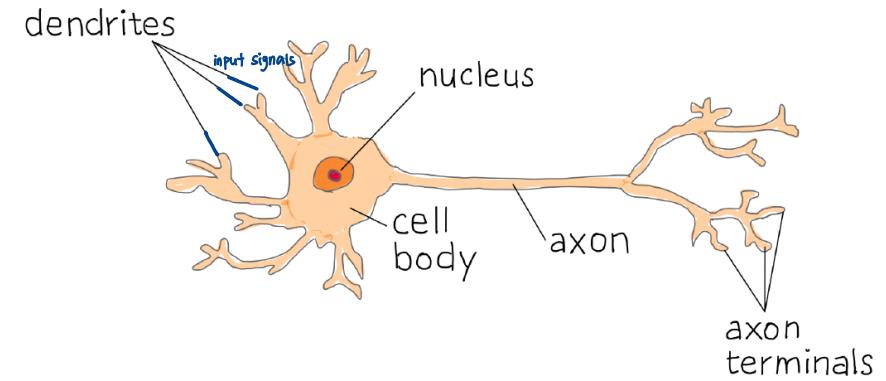
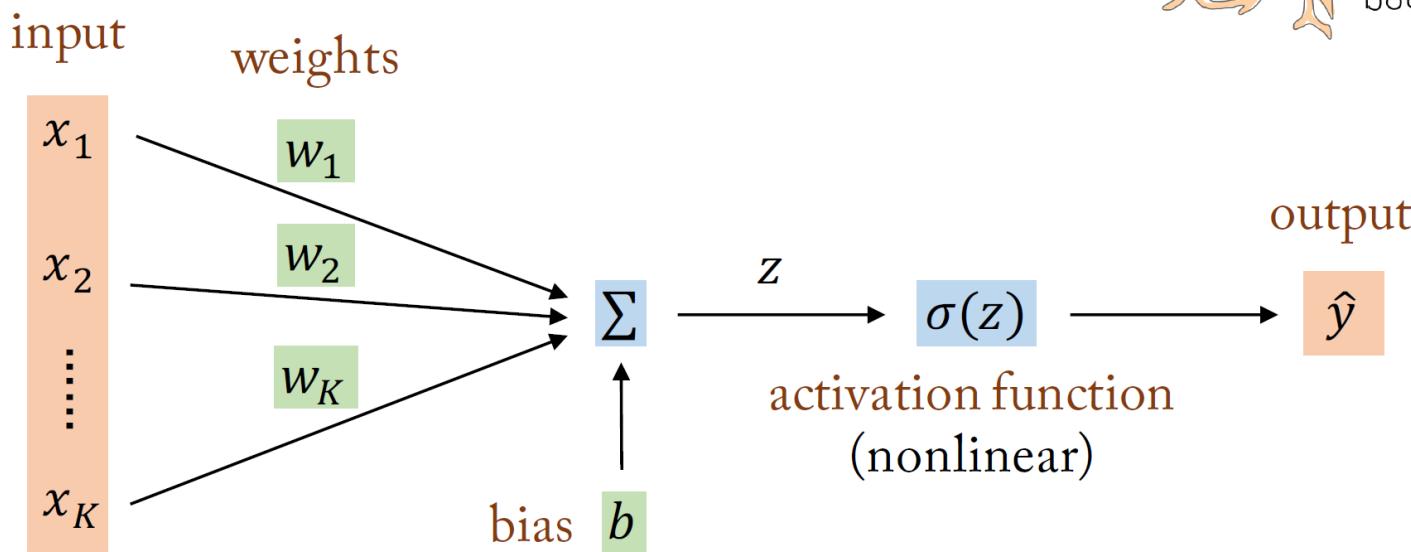




In Deep Learning, the function f is represented by a Neural Network.

Perceptron (neuron) [Frank Rosenblatt, 1957]

- The simplest structure of DL is one-layer model.



$$f : \mathbb{R}^K \rightarrow \mathbb{R}$$

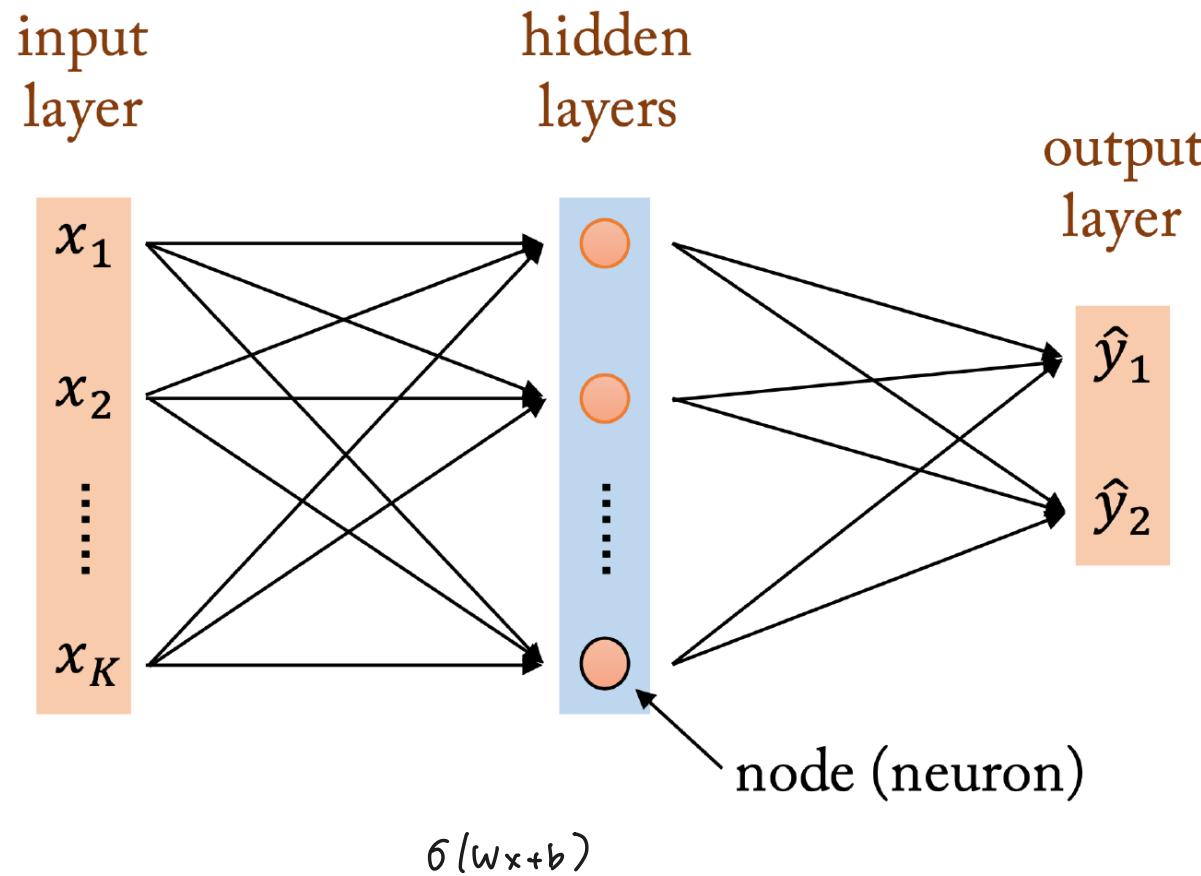
$$\hat{y} = \sigma(z) = \sigma(w_1x_1 + w_2x_2 + \dots + w_Kx_K + b)$$

\uparrow nonlinear \uparrow weighted sum (linear)

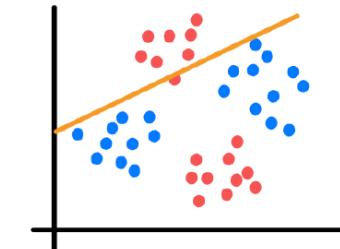
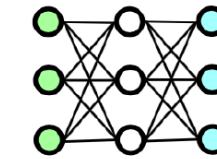
$= \sigma(\underbrace{w_x + b}_{\substack{\text{PK: } K \times 1 \\ \text{PK: } 1 \times K \\ \text{PK: Vector}}})$
similar to dense layer.

Multilayer Perceptron (MLP)

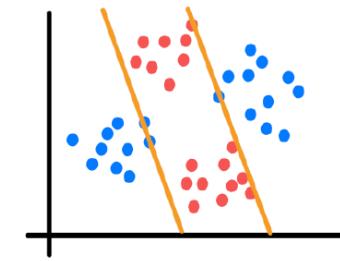
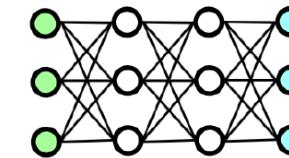
- Applying perceptron model layer by layer (= multi layer)



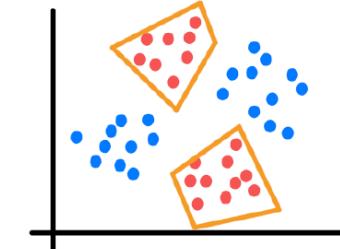
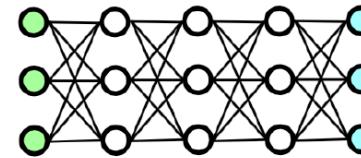
single layer



two layer



three layer

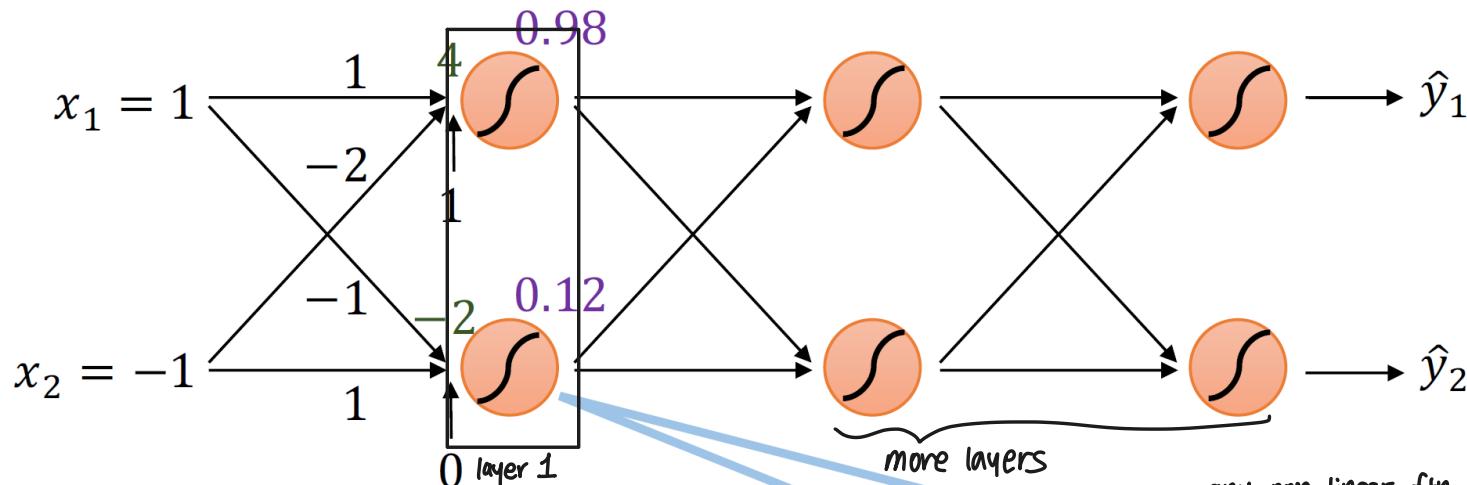


Forward pass computation

순전파

$$\sigma(w_1x_1 + w_2x_2 + \dots + w_Kx_K + b)$$

$$0.98 = \sigma(1 \cdot 1 + (-1) \cdot (-2) + 1)$$

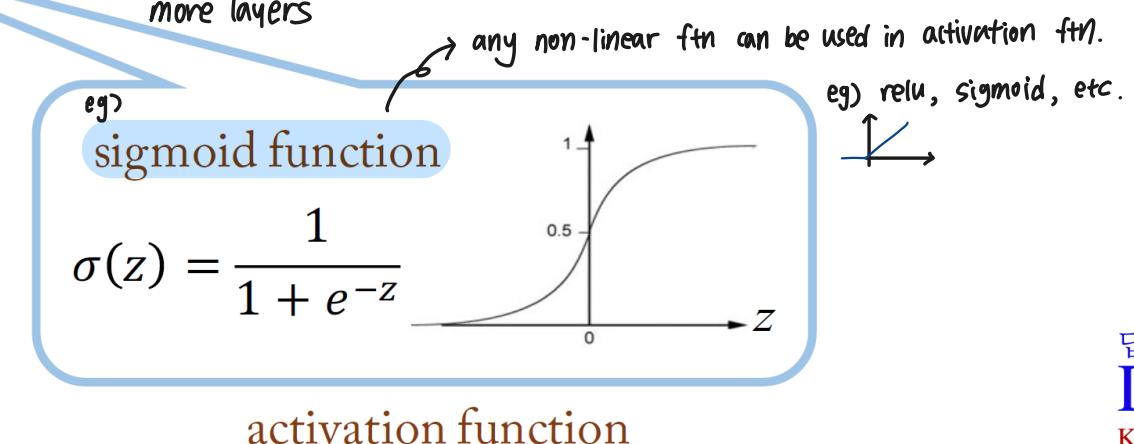


w, b trainable parameters.

After setting up DL model,

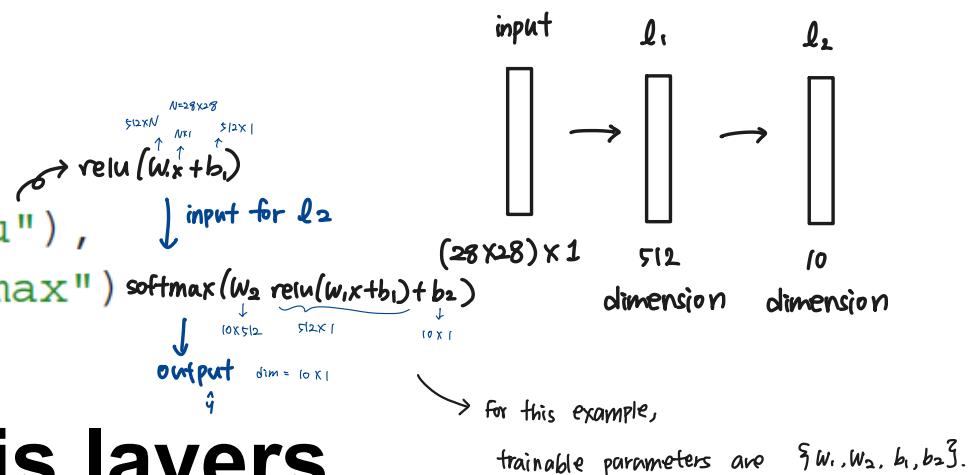
we should optimize weights (w) & bias (b)
to make model better.

$$\begin{aligned} \text{activation ftn} \quad & \sigma \left(\begin{bmatrix} w \\ b \end{bmatrix} \begin{bmatrix} x \\ 1 \end{bmatrix} + \begin{bmatrix} b \\ 0 \end{bmatrix} \right) \\ & = \sigma \left(\begin{bmatrix} 4 \\ -2 \end{bmatrix} \right) = \begin{bmatrix} 0.98 \\ 0.12 \end{bmatrix} \end{aligned}$$



Listing 2.2 The network architecture

```
from tensorflow import keras
from tensorflow.keras import layers
model = keras.Sequential([
    layers.Dense(512, activation="relu"),
    layers.Dense(10, activation="softmax")
])
```



Core building block of NN is layers

Layers extract *representations*

This step define our model $f(x; \text{weights})$



chapter02_mathematical-building-blocks.i

File Edit View Insert Runtime Tools Help Cannot save changes

 Share

The mathematical building blocks of neural networks

A first look at a neural network

- Data representations for neural networks
 - Scalars (rank-0 tensors)
 - Vectors (rank-1 tensors)
 - Matrices (rank-2 tensors)
 - Rank-3 and higher-rank tensors
 - Key attributes
 - Manipulating tensors in NumPy
 - The notion of data batches
 - Real-world examples of data tensors
 - Vector data
 - Timeseries data or sequence data
 - Image data
 - Video data
- The gears of neural networks: tensor

+ Code + Text |  Copy to Drive

✓ [8] train_images.shape
0s
(60000, 28, 28)

✓ [9] len(train_labels)
0s
60000

✓ [10]  train_labels
0s
 array([5, 0, 4, ..., 5, 6, 8], dtype=uint8)

✓ [11] test_images.shape
0s
(10000, 28, 28)

✓ [12] len(test_labels)
0s
10000

✓ [13] test_labels
0s
 array([7, 2, 1, ..., 4, 5, 6], dtype=uint8)

The network architecture

Warning: you are connected to a GPU runtime, but not utilizing the GPU. [Change to a standard runtime](#)

– 83 – completed at 12:36 PM

Step 1) define our model $f_\theta(x)=y$

Step 2) define criteria \rightarrow loss function

How the network measure its performance.

What's next?

Loss function: How the network measure its performance

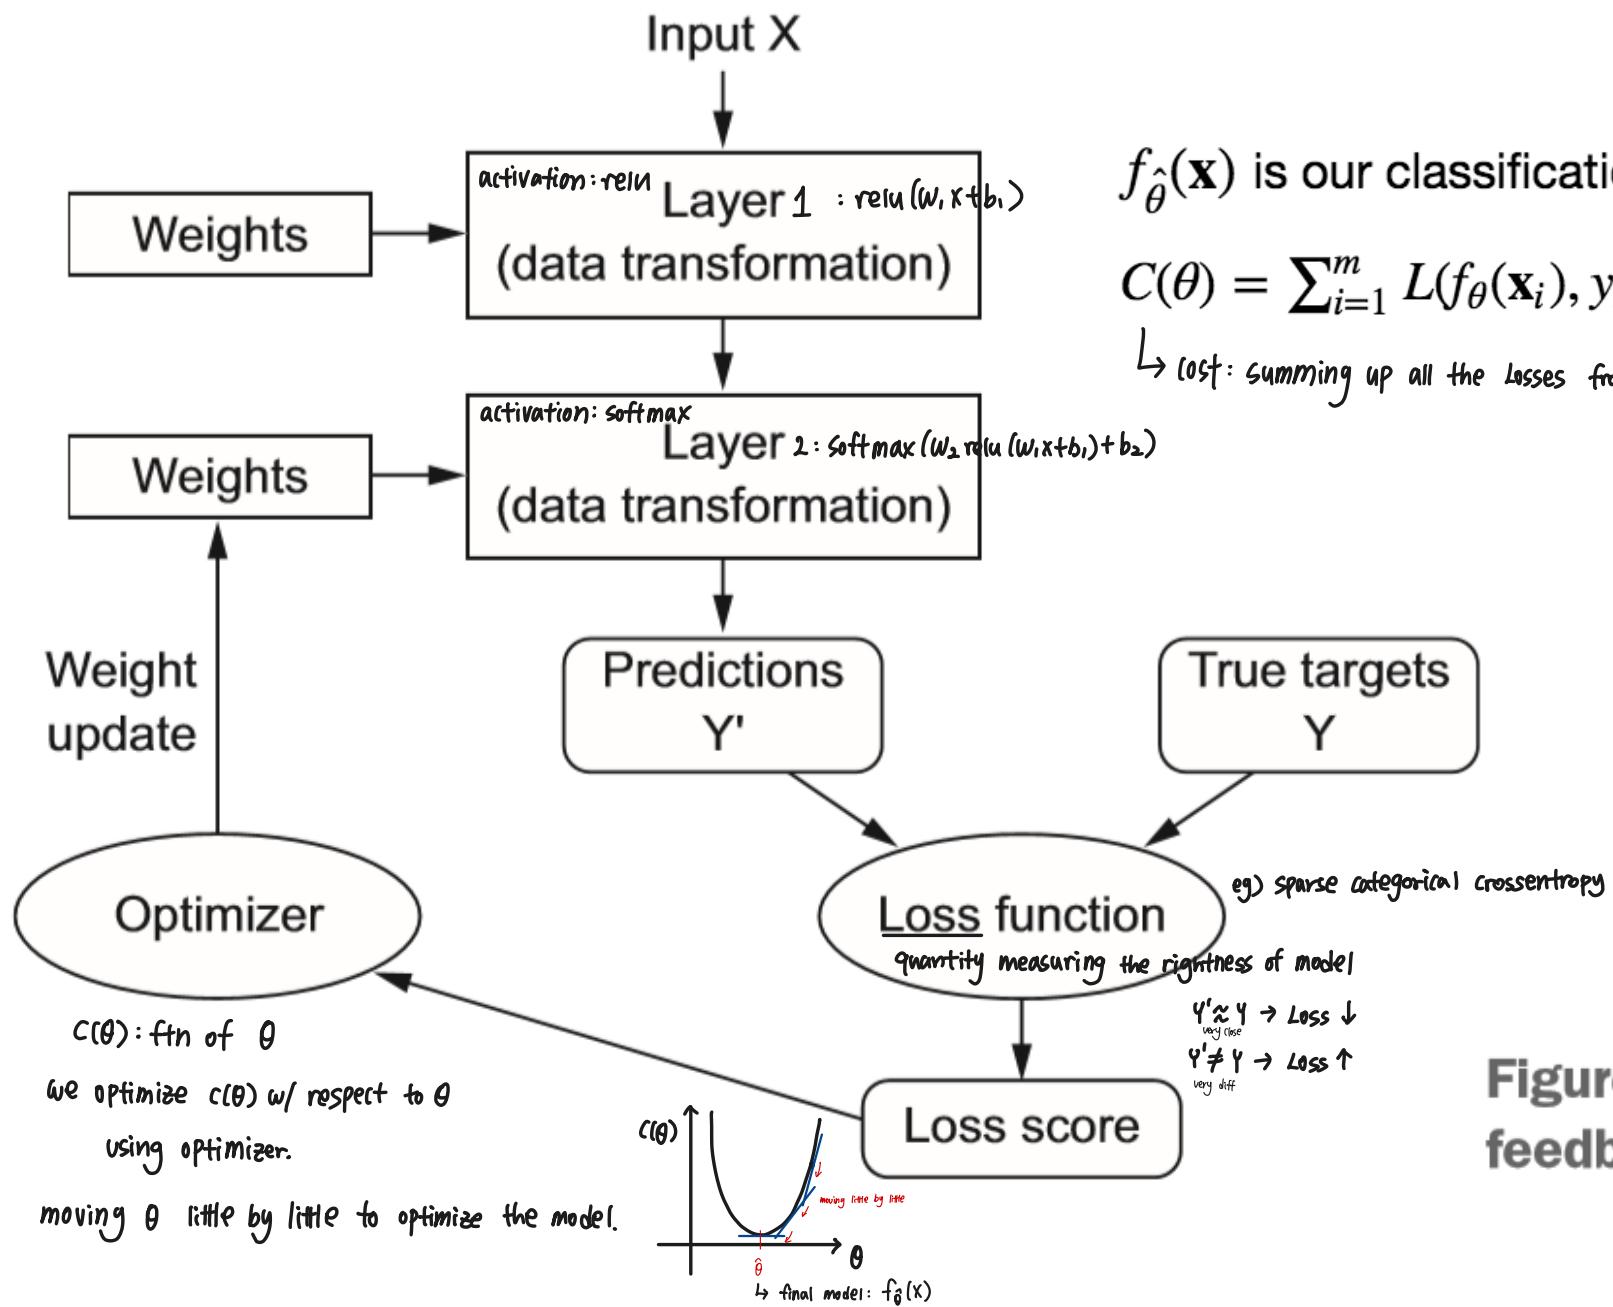
Optimizer: Mechanism through which the network will update itself based on the data it sees and its loss function

Metrics to monitor during training and testing: Accuracy

Listing 2.3 The compilation step

```
model.compile(optimizer="rmsprop",
              loss="sparse_categorical_crossentropy",
              metrics=["accuracy"] )
```





$\hat{f}_{\hat{\theta}}(\mathbf{x})$ is our classification model with estimated parameters $\hat{\theta}$

$$C(\theta) = \sum_{i=1}^m L(f_\theta(\mathbf{x}_i), y_i)$$

→ cost: summing up all the losses from $\underbrace{i \text{ to } m}_{\text{all indicies.}} / \# \text{ training data set}$

all indicies. / # training data set

True targets

e.g.) sparse categorical crossentropy

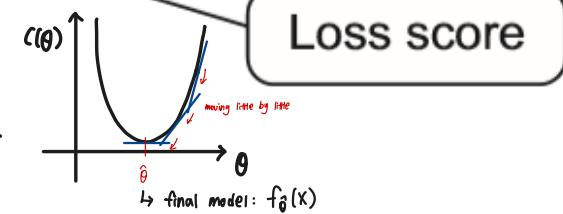
Quantity measuring the rightness of model

Quantity measuring the rightness of model

$\gamma' \approx \gamma \rightarrow \text{Loss} \downarrow$
very close

$\gamma' \neq \gamma \rightarrow \text{Loss} \uparrow$
very diff

Figure 1.9 The loss score is used as a feedback signal to adjust the weights.





chapter02_mathematical-building-blocks.i

File Edit View Insert Runtime Tools Help Cannot save changes

 Share  

 RAM  Disk

Editing

^

 Table of contents 

The mathematical building blocks of neural networks

A first look at a neural network

Data representations for neural networks

Scalars (rank-0 tensors)

Vectors (rank-1 tensors)

Matrices (rank-2 tensors)

Rank-3 and higher-rank tensors

Key attributes

Manipulating tensors in NumPy

The notion of data batches

Real-world examples of data tensors

Vector data

Timeseries data or sequence data

Image data

Video data

The gears of neural networks: tensor operations

Element-wise operations

 + Code  + Text 

```
[14] from tensorflow import keras
from tensorflow.keras import layers
model = keras.Sequential([
    layers.Dense(512, activation="relu"),
    layers.Dense(10, activation="softmax")
])
```

The compilation step

 + Code  + Text       

```
model.compile(optimizer="rmsprop",
              loss="sparse_categorical_crossentropy",
              metrics=["accuracy"])
```

```
[ ]
```

Preparing the image data

```
[ ] train_images = train_images.reshape((60000, 28 * 28))
train_images = train_images.astype("float32") / 255
test_images = test_images.reshape((10000, 28 * 28))
test_images = test_images.astype("float32") / 255
```

"Fitting" the model

```
[ ] model.fit(train_images, train_labels, epochs=5, batch_size=128)
```

 0s completed at 1:03 PM

오후 1:04

2022-03-06

 검색

Rescale values from [0,255] to [0,1]

Listing 2.4 Preparing the image data

```
train_images = train_images.reshape((60000, 28 * 28))  
train_images = train_images.astype("float32") / 255  
test_images = test_images.reshape((10000, 28 * 28))  
test_images = test_images.astype("float32") / 255
```

Fitting the model

Listing 2.5 “Fitting” the model

```
>>> model.fit(train_images, train_labels, epochs=5, batch_size=128)  
Epoch 1/5  
60000/60000 [=====] - 5s - loss: 0.2524 - acc: 0.9273  
Epoch 2/5  
51328/60000 [=====>.....] - ETA: 1s - loss: 0.1035 - acc: 0.9692
```

using training data 5 times
every iteration of the training
how many mini batch sample that we are using?



Make Predictions

Listing 2.6 Using the model to make predictions

```
>>> test_digits = test_images[0:10]
>>> predictions = model.predict(test_digits)
>>> predictions[0]
array([1.0726176e-10, 1.6918376e-10, 6.1314843e-08, 8.4106023e-06,
       2.9967067e-11, 3.0331331e-09, 8.3651971e-14, 9.9999106e-01,
       2.6657624e-08, 3.8127661e-07], dtype=float32)
```

the result would be "7".

Evaluating the model

Listing 2.7 Evaluating the model on new data

```
>>> test_loss, test_acc = model.evaluate(test_images, test_labels)
>>> print(f"test_acc: {test_acc}")
test_acc: 0.9785
```



Table of contents

- The mathematical building blocks of neural networks
- A first look at a neural network**
- Data representations for neural networks
 - Scalars (rank-0 tensors)
 - Vectors (rank-1 tensors)
 - Matrices (rank-2 tensors)
 - Rank-3 and higher-rank tensors
 - Key attributes
 - Manipulating tensors in NumPy
 - The notion of data batches
 - Real-world examples of data tensors
 - Vector data
 - Timeseries data or sequence data
 - Image data
 - Video data
- The gears of neural networks: tensor operations
- Element-wise operations

+ Code + Text 

3s  from tensorflow import keras
from tensorflow.keras import layers
model = keras.Sequential([
 layers.Dense(512, activation="relu"),
 layers.Dense(10, activation="softmax")
])

The compilation step

[22] model.compile(optimizer="rmsprop",
 loss="sparse_categorical_crossentropy",
 metrics=["accuracy"])

 |

Preparing the image data

[] train_images = train_images.reshape((60000, 28 * 28))
train_images = train_images.astype("float32") / 255
test_images = test_images.reshape((10000, 28 * 28))
test_images = test_images.astype("float32") / 255

"Fitting" the model

[] model.fit(train_images, train_labels, epochs=5, batch_size=128)

✓ 0s completed at 1:04 PM

CH2. Mathematical Building Blocks of Deep-learning (2)

Il-Youp Kwak
Chung-Ang University



Data representations for neural networks

Tensor is a container for data

0D tensors (Scalars)

```
>>> import numpy as np  
>>> x = np.array(12)  
>>> x  
array(12)  
>>> x.ndim  
0
```

1D tensors (Vectors)

```
>>> x = np.array([12, 3, 6, 14])  
>>> x  
array([12, 3, 6, 14])  
>>> x.ndim  
1
```

2D tensors (Matrices)

```
>>> x = np.array([[5, 78, 2, 34, 0],  
[6, 79, 3, 35, 1],  
[7, 80, 4, 36, 2]])  
>>> x.ndim  
2
```



3D tensors and higher-dimensional tensors

```
>>> x = np.array([[[5, 78, 2, 34, 0],  
[6, 79, 3, 35, 1],  
[7, 80, 4, 36, 2]],  
[[5, 78, 2, 34, 0],  
[6, 79, 3, 35, 1],  
[7, 80, 4, 36, 2]],  
[[5, 78, 2, 34, 0],  
[6, 79, 3, 35, 1],  
[7, 80, 4, 36, 2]]])  
>>> x.ndim  
3
```



Key attributes for *tensors*

Number of axes (rank): can be checked by `.ndim` for numpy objects

0D, 1D, 2D
rank

Shape: tuple of integers that describes how many dimensions the tensor has along each axis. Can be checked by `.shape` for numpy objects

Data type (dtype): type of the data in tensor. Ex: float32, uint8, float64, and so on.



Key attributes for *tensors*

```
from keras.datasets import mnist  
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()  
  
>>> print(train_images.ndim)  
3  
  
>>> print(train_images.shape)  
(60000, 28, 28)  
  
>>> print(train_images.dtype)  
uint8
```



Displaying Digits

```
digit = train_images[4]

import matplotlib.pyplot as plt
plt.imshow(digit, cmap=plt.cm.binary)
plt.show()
```

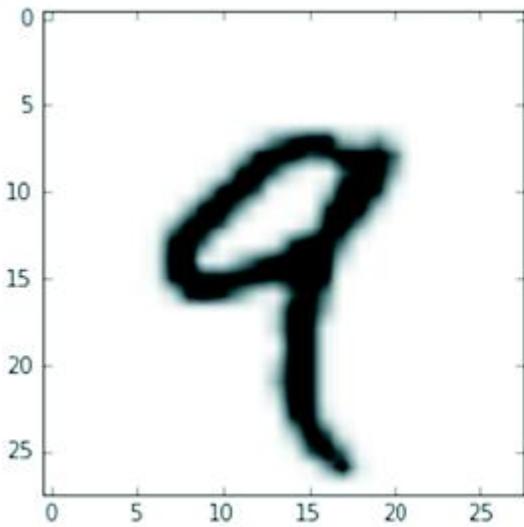


Figure 2.2 The fourth sample in our dataset



news study data science 학회사이트 업무관련 기타 Bookmark Manager Gmail YouTube 지도 뉴스 번역 SR Translate scm Data 기타 북마크 읽기 목록

chapter02_mathematical-building-blocks.i

File Edit View Insert Runtime Tools Help Cannot save changes

Table of contents

- The mathematical building blocks of neural networks
- A first look at a neural network
- Data representations for neural networks**
 - Scalars (rank-0 tensors)
 - Vectors (rank-1 tensors)
 - Matrices (rank-2 tensors)
 - Rank-3 and higher-rank tensors
 - Key attributes
 - Manipulating tensors in NumPy
 - The notion of data batches
 - Real-world examples of data tensors
 - Vector data
 - Timeseries data or sequence data
 - Image data
 - Video data
 - The gears of neural networks: tensor operations
 - Element-wise operations

+ Code + Text Copy to Drive x.ndim

Key attributes

```
[ ] from tensorflow.keras.datasets import mnist  
[ ] (train_images, train_labels), (test_images, test_labels) = mnist.load_data()  
[ ] train_images.ndim  
[ ] train_images.shape  
[ ] train_images.dtype
```

Displaying the fourth digit

```
[ ] import matplotlib.pyplot as plt  
[ ] digit = train_images[4]  
[ ] plt.imshow(digit, cmap=plt.cm.binary)  
[ ] plt.show()  
[ ] train_labels[4]
```

Manipulating tensors in NumPy

✓ 1s completed at 1:16 PM

Manipulating tensors in Numpy

Tensor slicing

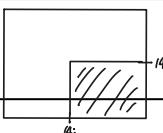
```
>>> my_slice = train_images[10:100]
>>> print(my_slice.shape)           ||
(90, 28, 28)
>>> my_slice = train_images[10:100, :, :]
>>> my_slice.shape                ||
(90, 28, 28)                    || same
>>> my_slice = train_images[10:100, 0:28, 0:28]
>>> my_slice.shape
(90, 28, 28)
```



Manipulating tensors in Numpy

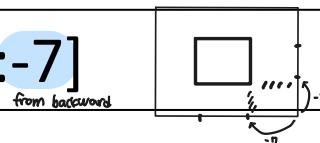
Select 14×14 pixels in bottom-right corner of all images

```
my_slice = train_images[:, 14:, 14:]
```



Crop 14×14 pixels centered in the middle

```
my_slice = train_images[:, 7:-7, 7:-7]
```



```
plt.imshow(my_slice[4], cmap=plt.cm.binary)  
plt.show()
```





chapter02_mathematical-building-blocks.i

File Edit View Insert Runtime Tools Help Cannot save changes

 Share

Table of contents	
The mathematical building blocks of neural networks	
A first look at a neural network	
Data representations for neural networks	
Scalars (rank-0 tensors)	
Vectors (rank-1 tensors)	
Matrices (rank-2 tensors)	
Rank-3 and higher-rank tensors	
Key attributes	
Manipulating tensors in NumPy	
The notion of data batches	
Real-world examples of data tensors	
Vector data	
Timeseries data or sequence data	
Image data	
Video data	
The gears of neural networks: tensor operations	
Element-wise operations	

 + Code  + Text 

[47] train_labels[4]

9

 Manipulating tensors in NumPymy_slice  Loading... train_images[10:100]

my_slice.shape

[] my_slice = train_images[10:100, :, :]
my_slice.shape[] my_slice = train_images[10:100, 0:28, 0:28]
my_slice.shape

[] my_slice = train_images[:, 14:, 14:]

[] my_slice = train_images[:, 7:-7, 7:-7]

 The notion of data batches

[] batch = train_images[:128]

[] batch = train_images[128:256]

0s completed at 1:30 PM

The notion of data batches

The 1st axis of data tensors in deep-learning will be sample axis 60,000

Deep-learning models don't process an entire dataset at once

process little by little (using batches)

Like

mini batch gradient descent algorithm

```
# 1st batch  
batch = train_images[:128]  
# 2nd batch  
batch = train_images[128:256]  
# nth batch  
batch = train_images[128 * n:128 * (n + 1)]
```



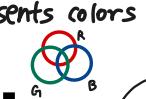
Real-world examples of data tensors

Vector data – 2D tensor of shape (samples, features)

Timeseries data or sequence data – 3D tensors of shape (samples, timesteps, features)

(cf. MNIST : 3D :: black-white image \Rightarrow it means 1 channel)

Images – 4D tensors of shape (samples, height, width, channels) **or** (samples, channels, height, width)




Video – 5D tensors of shape (samples, frames, height, width, channels) **or** (samples, frames, channels, height, width)



Vector Data

Vector data – 2D tensor of shape (samples, features)

Example:

Dataset of people, where we consider each person's age, ZIP code, and *income*.

Each person can be characterized as a vector of 3 values
Entire dataset of 100,000 people can be stored in a 2D tensor of shape (100000, 3).



Timeseries data or sequence data

Timeseries data or sequence data – 3D tensors of shape (samples, timesteps, features)

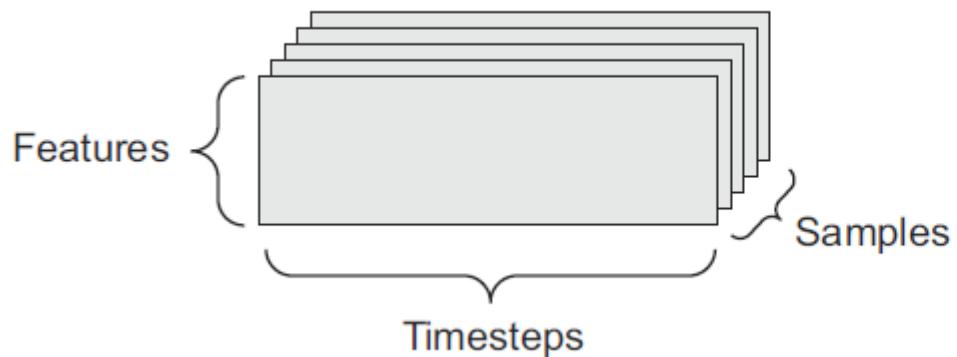


Figure 2.3 A 3D timeseries data tensor

Dataset of stock prices. Every minute, we store the *current price, the highest and the lowest price in the past minute*. Encoded as a 2D tensor of shape $(390, 3)$ for 390 minutes in a trading day
250 days' data stored in 3D tensor of shape $(250, 390, 3)$.



Image data

Images – 4D tensors of shape (samples, height, width, channels) or (samples, channels, height, width)

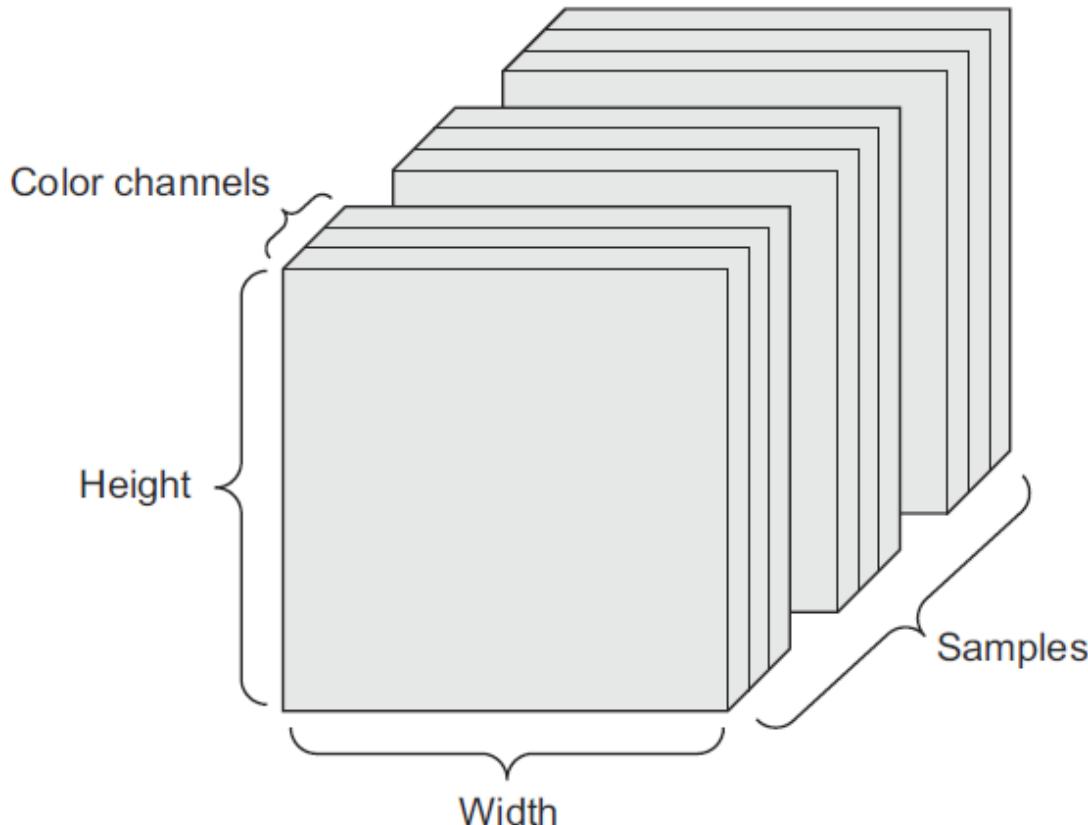


Figure 2.4 A 4D image data tensor (channels-first convention)



Video data

Video – 5D tensors of shape (samples, frames, height, width, channels) or (samples, frames, channels, height, width)

Example:

A 60-second, 144×256 YouTube video clip sampled at 4 frames per second would have 240 frames. A batch of four such video clips would be stored in a tensor of shape (4, 240, 144, 256, 3)



The **gears** of neural networks: **tensor operations**

Tensor operations applied to tensors of numeric data.

Example:

```
keras.layers.Dense(512, activation='relu')
```

Unpack this

output = relu($\underset{\text{max}(x, 0)}{\text{dot}}(\underset{\text{dot product}}{\text{dot}}(\underset{512 \times 700}{W}, \underset{\text{eg}}{\text{input}}) \underset{\text{addition}}{+} \underset{700 \times 1}{b})$)

Three tensor operations:

dot product (dot) between input tensor and a tensor W

addition (+) between the resulting 2D tensor and vector b

relu operation: $\text{relu}(x)$ is $\max(x, 0)$.



Element-wise operations

Ex) relu and addition

eg) $(1, 2, 3) + (0, 0, 1) = (1, 2, 4)$
Calculate elementwise!

Naïve implementation of an element-wise relu operation:

X should be matrix form.

```
def naive_relu(x):
    assert len(x.shape) == 2, 'x.ndim() must 2'
    x = x.copy()
    for i in range(x.shape[0]):
        for j in range(x.shape[1]):
            x[i, j] = max(x[i, j], 0)
    return x
```



← → C ⌘ colab.research.google.com/github/fchollet/deep-learning-with-python-notebooks/blob/master/chapter02_mathematical-building-blocks.ipynb#scrollTo... ☆ EX | 업데이트 :

news study data science 학회사이트 업무관련 기타 Bookmark Manager Gmail YouTube 지도 뉴스 번역 SR Translate scm Data » 기타 북마크 읽기 목록

CO chapter02_mathematical-building-blocks.i

File Edit View Insert Runtime Tools Help Cannot save changes

Table of contents

The mathematical building blocks of neural networks

A first look at a neural network

Data representations for neural networks

Scalars (rank-0 tensors)

Vectors (rank-1 tensors)

Matrices (rank-2 tensors)

Rank-3 and higher-rank tensors

Key attributes

Manipulating tensors in NumPy

The notion of data batches

Real-world examples of data tensors

Vector data

Timeseries data or sequence data

Image data

Video data

The gears of neural networks: tensor operations

Element-wise operations

+ Code + Text Copy to Drive

RAM Disk

Editing

12

0 2 4 6 8 10 12

▼ The notion of data batches

```
[ ] batch = train_images[:128]
```

```
[ ] batch = train_images[128:256]
```

```
[ ] n = 3  
batch = train_images[128 * n:128 * (n + 1)]
```

Real-world examples of data tensors

Vector data

Timeseries data or sequence data

Image data

✓ 0s completed at 1:36 PM

오늘 1:59 2022-03-06

Element-wise operations

Ex) relu and addition

Naïve implementation of addition:

```
def naive_add(x, y):
    assert len(x.shape) == 2
    assert x.shape == y.shape

    x = x.copy() ## avoid overwriting
    for i in range(x.shape[0]):
        for j in range(x.shape[1]):
            x[i, j] += y[i, j]
    return x
```

input 으로 넣은 x가 그대로 살아남기
and original values changing



Element-wise operations

Element-wise relu:

```
import numpy as np  
z = np.maximum(z, 0.)
```

Element-wise add:

```
z = x + y
```



Broadcasting

What happens with addition when the shapes of the two tensors being added differ?

Broadcasting consists of two steps:

1. Axes (called broadcast axes) are added to the smaller tensor to match the `ndim` of the larger tensor.
2. The smaller tensor is repeated alongside these new axes to match the full shape of the larger tensor.

eg) $\underbrace{X \& Y}$
dims are diff
 $X + Y$
 $X.ndim > Y.ndim \Rightarrow Y \text{ Broadcasting}$



Broadcasting

Example:

```
import numpy as np
x = np.array([1,2,3,4,5]) ①
y = np.array([[1,1,1,1,1], [1,1,1,1,1]]) ②
```

$$\begin{matrix} \text{Broadcast} \leftarrow & \begin{bmatrix} 12345 \\ 12345 \\ 11111 \end{bmatrix} + \begin{bmatrix} 11111 \\ 11111 \end{bmatrix} \\ & = \begin{bmatrix} 23456 \\ 23456 \end{bmatrix} \end{matrix}$$

```
def naive_add_matrix_and_vector(x, y):
    assert len(x.shape) == 2
    assert len(y.shape) == 1
    assert x.shape[1] == y.shape[0]

    x = x.copy()
    for i in range(x.shape[0]):
        for j in range(x.shape[1]):
            x[i, j] += y[j]
    return x
```

x is a **2D Numpy tensor**.

y is a **Numpy vector**.

Avoid overwriting the input tensor.



Broadcasting

if one tensor has shape $(a, b, \dots n, n + 1, \dots m)$ and the other has shape $(n, n + 1, \dots m)$. The broadcasting will then automatically happen for axes a through $n - 1$.

```
import numpy as np  
  
x = np.random.random((64, 3, 32, 10))  
y = np.random.random((32, 10))  
z = np.maximum(x, y)
```

x is a random tensor with shape (64, 3, 32, 10).

y is a random tensor with shape (32, 10).

The output z has shape (64, 3, 32, 10) like x.



news study data science 학회사이트 업무관련 기타 Bookmark Manager Gmail YouTube 지도 뉴스 번역 SR Translate scm Data 기타 북마크 읽기 목록

chapter02_mathematical-building-blocks.i

File Edit View Insert Runtime Tools Help Cannot save changes

Table of contents

- The notion of data batches
- Real-world examples of data tensors
- Vector data
- Timeseries data or sequence data
- Image data
- Video data
- The gears of neural networks: tensor operations
- Element-wise operations**
- Broadcasting
- Tensor product
- Tensor reshaping
- Geometric interpretation of tensor operations
- A geometric interpretation of deep learning
- The engine of neural networks: gradient-based optimization
- What's a derivative?
- Derivative of a tensor operation: the gradient

+ Code + Text Copy to Drive RAM Disk Editing

```
[73] x1
0s
array([[1, 0, 3],
       [2, 3, 0]])
```

```
[ ] def naive_add(x, y):
    assert len(x.shape) == 2
    assert x.shape == y.shape
    x = x.copy()
    for i in range(x.shape[0]):
        for j in range(x.shape[1]):
            x[i, j] += y[i, j]
    return x
```

```
[ ] import time

x = np.random.random((20, 100))
y = np.random.random((20, 100))

t0 = time.time()
for _ in range(1000):
    z = x + y
    z = np.maximum(z, 0.)
print("Took: {:.2f} s".format(time.time() - t0))
```

```
[ ] t0 = time.time()
for _ in range(1000):
```

0s completed at 2:16 PM

Tensor dot

Dot product in math

$$(a, b, c) \cdot (d, e, f) = ad + be + cf$$

```
import numpy as np  
z = np.dot(x, y)
```

```
def naive_vector_dot(x, y):  
    assert len(x.shape) == 1  
    assert len(y.shape) == 1  
    assert x.shape[0] == y.shape[0]  
    z = 0.  
    for i in range(x.shape[0]):  
        z += x[i] * y[i]  
    return z
```

x and y are NumPy vectors.



Matrix vector dot

```
def naive_matrix_vector_dot(x, y):  
    assert len(x.shape) == 2  
    assert len(y.shape) == 1  
    assert x.shape[1] == y.shape[0]  
    z = np.zeros(x.shape[0])  
    for i in range(x.shape[0]):  
        for j in range(x.shape[1]):  
            z[i] += x[i, j] * y[j]  
    return z
```

$$\begin{matrix} w_x \\ a \times b \quad b \times 1 \end{matrix}$$

x is a NumPy matrix. (w)

y is a NumPy vector. (x)

The first dimension of x
must be the same as the
0th dimension of y!

This operation returns a vector
of 0s with the same shape as y.



Matrix vector dot

```
def naive_matrix_vector_dot(x, y):
    assert len(x.shape) == 2
    assert len(y.shape) == 1
    assert x.shape[1] == y.shape[0]
    z = np.zeros(x.shape[0])
    for i in range(x.shape[0]):
        for j in range(x.shape[1]):
            z[i] += x[i, j] * y[j]
    return z
```

x is a NumPy matrix.

y is a NumPy vector.

The first dimension of x must be the same as the 0th dimension of y!

This operation returns a vector of 0s with the same shape as y.

Could also use the previous codes

```
def naive_matrix_vector_dot(x, y):
    z = np.zeros(x.shape[0])
    for i in range(x.shape[0]):
        z[i] = naive_vector_dot(x[i, :], y)
    return z
```

$$\begin{array}{c} x \\ \hline i \end{array}
 \begin{array}{c} y \\ \hline \end{array}
 =
 \begin{array}{c} z \\ \hline \text{ith} \end{array}$$



Dot product for 2D tensors

This operation returns a matrix of 0s with a specific shape.

```
def naive_matrix_dot(x, y):
    assert len(x.shape) == 2
    assert len(y.shape) == 2
    assert x.shape[1] == y.shape[0]
    z = np.zeros((x.shape[0], y.shape[1]))
    for i in range(x.shape[0]):
        for j in range(y.shape[1]):
            row_x = x[i, :]
            column_y = y[:, j]
            z[i, j] = naive_vector_dot(row_x, column_y)
    return z
```

x and y are NumPy matrices.

The first dimension of x must be the same as the 0th dimension of y!

Iterates over the rows of x ...
... and over the columns of y.

The gears of neural networks: Tensor operations

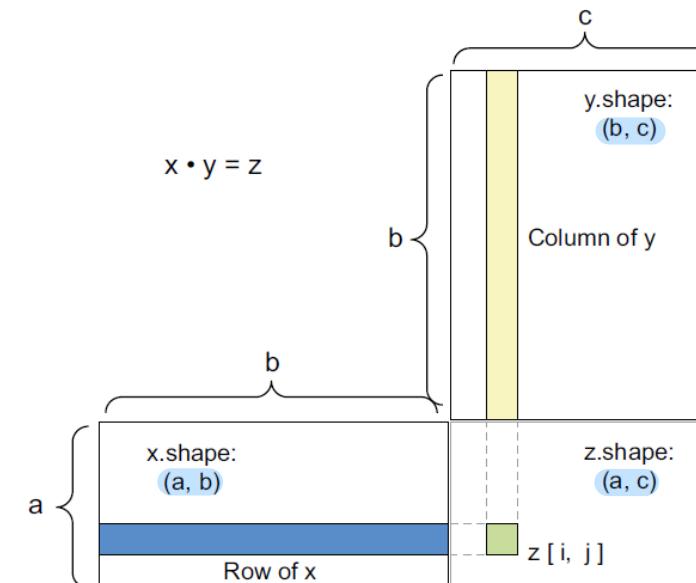


Figure 2.5 Matrix dot-product box diagram



Tensor dot

More generally, you can take the dot product between higher-dimensional tensors, following the same rules for shape compatibility as outlined earlier for the 2D case:

$$\begin{array}{c} a \times b \cdot b \times c \Rightarrow a \times c \\ (a, b, c, d) . (d,) \xrightarrow{1D} (a, b, c) \\ (a, b, c, d) . (d, e) \xrightarrow{1D} (a, b, c, e) \end{array}$$



Tensor reshaping

we preprocessed the digits data before feeding it into our network:

```
train_images = train_images.reshape((60000, 28 * 28))
```

Try it:

```
x = np.array([[0., 1.], [2., 3.], [4., 5.]])
```

```
x = x.reshape((6, 1))
```

```
x = x.reshape((2, 3))
```



chapter02_mathematical-building-blocks.i

File Edit View Insert Runtime Tools Help Cannot save changes

Table of contents

- The notion of data batches
- Real-world examples of data tensors
- Vector data
- Timeseries data or sequence data
- Image data
- Video data
- The gears of neural networks: tensor operations
 - Element-wise operations
- Broadcasting**
 - Tensor product
 - Tensor reshaping
 - Geometric interpretation of tensor operations
 - A geometric interpretation of deep learning
- The engine of neural networks: gradient-based optimization
 - What's a derivative?
 - Derivative of a tensor operation: the gradient

+ Code + Text Copy to Drive RAM Disk Editing

```
[ ] import numpy as np  
x = np.random.random((64, 3, 32, 10))  
y = np.random.random((32, 10))  
z = np.maximum(x, y)
```

Tensor product

```
[ ] x = np.random.random((32,))  
y = np.random.random((32,))  
z = np.dot(x, y)
```

```
[ ] def naive_vector_dot(x, y):  
    assert len(x.shape) == 1  
    assert len(y.shape) == 1  
    assert x.shape[0] == y.shape[0]  
    z = 0.  
    for i in range(x.shape[0]):  
        z += x[i] * y[i]  
    return z
```

```
[ ] def naive_matrix_vector_dot(x, y):  
    assert len(x.shape) == 2  
    assert len(y.shape) == 1  
    assert x.shape[1] == y.shape[0]
```

0s completed at 2:19 PM

Geometric interpretation of tensor operations

Vector as an arrow

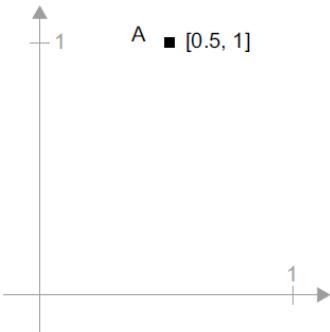


Figure 2.6 A point in a 2D space

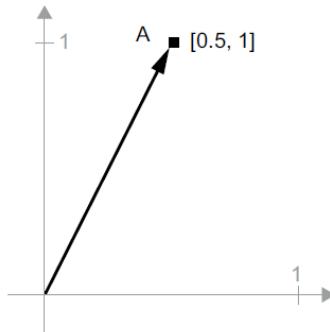


Figure 2.7 A point in a 2D space pictured as an arrow

Sum of two vectors

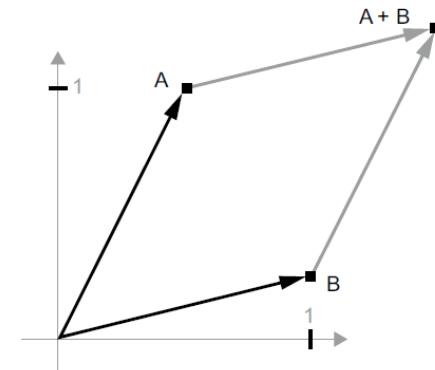


Figure 2.8 Geometric interpretation of the sum of two vectors

Translation as vector addition

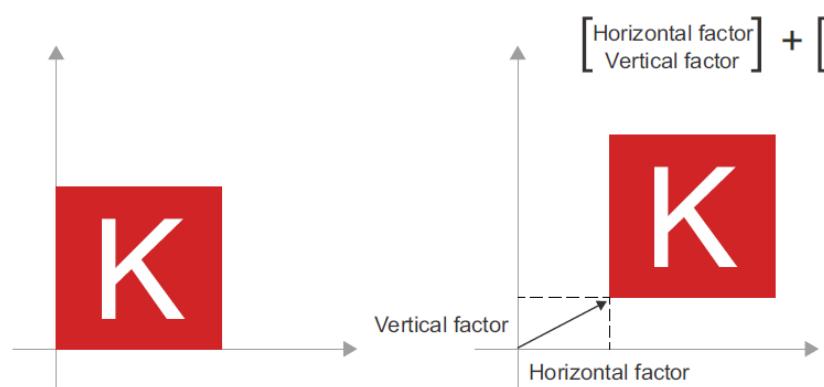


Figure 2.9 2D translation as a vector addition

2D rotation

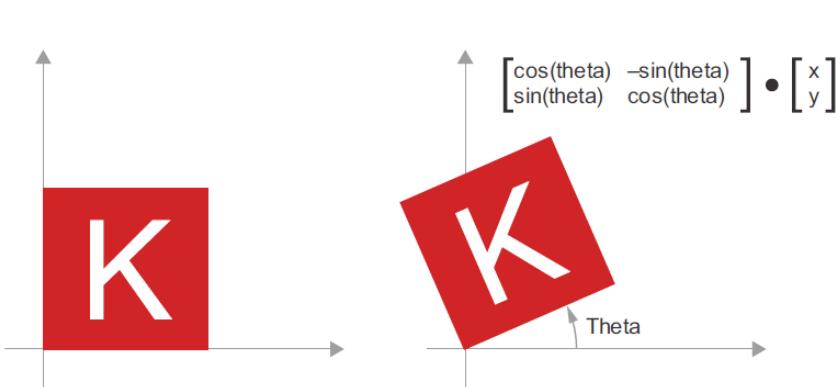
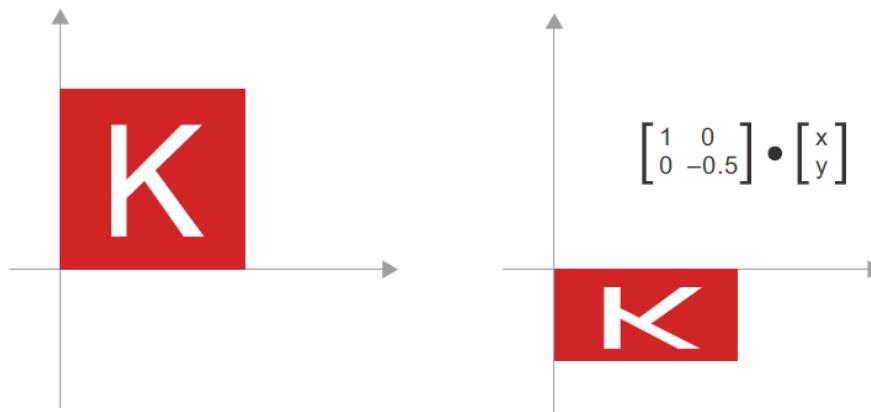


Figure 2.10 2D rotation (counterclockwise) as a dot product

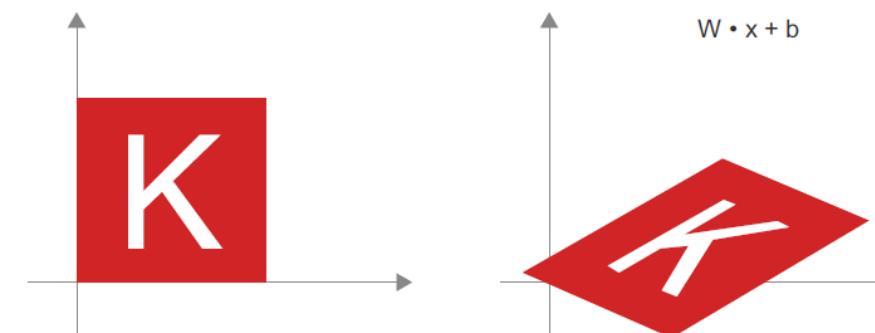


Geometric interpretation of tensor operations

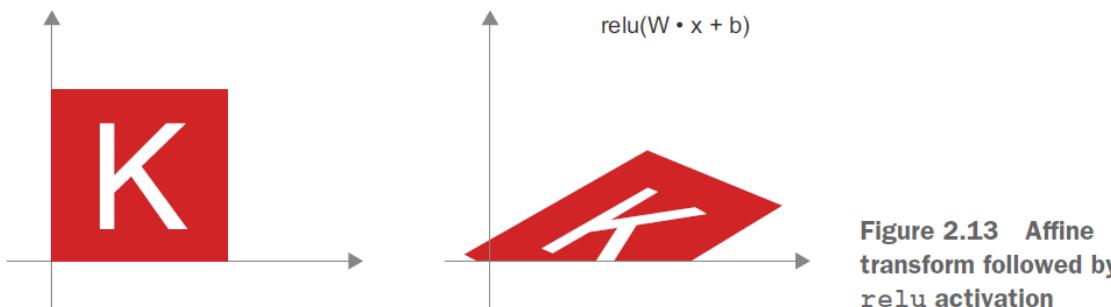
2D scaling



Affine transform



Affine transform followed by relu



Geometric interpretation of deep learning

Neural networks consist entirely of chains of tensor operations

These tensor operations are just simple geometric transformations of the input data

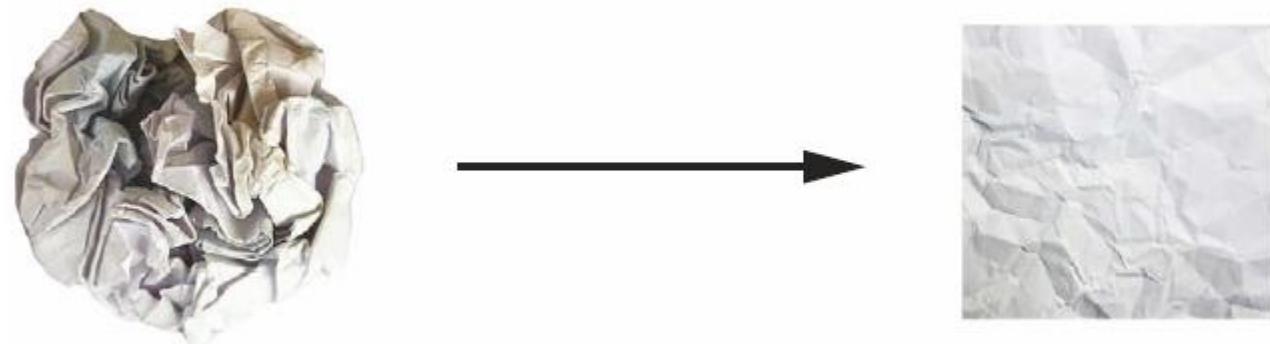


Figure 2.14 Uncrumpling a complicated manifold of data



Thank you! 😊

