

CH4. Getting Started with Neural Networks

Il-Youp Kwak
Chung-Ang University



Getting Started with Neural Networks

Common ML examples

1. **Classifying movie reviews: a binary classification example**
2. **Classifying newswires: a multiclass classification example**
3. **Predicting house prices: a regression example**



Classifying movie reviews: a binary classification example

The IMDB dataset:

- **50,000 highly polarized reviews**
- **25,000 reviews for training and 25,000 reviews for testing**
- **Each set consisting of 50% negative and 50% positive reviews**

데이터 균형 맞아야 함!



- **Read the data**

Listing 4.1 Loading the IMDB dataset

```
from tensorflow.keras.datasets import imdb
(train_data, train_labels), (test_data, test_labels) = imdb.load_data(
    num_words=10000)
    → this dataset is included in keras.datasets
    we'll use this number of frequently used words in review.
```

- **Decode a review**

Listing 4.2 Decoding reviews back to text

```
word_index = imdb.get_word_index() ← word_index is a dictionary mapping
reverse_word_index = dict( words to an integer index.
    [ (value, key) for (key, value) in word_index.items() ] ) ←
decoded_review = " ".join( [ reverse_word_index.get(i - 3, "?") for i in train_data[0] ] ) ←
    iterate ← Reverses it,
    this would be sequence of numbers ← mapping
    eg) [11, 3, 90, 101] ← integer indices
    to words
```

Decodes the review. Note that the indices are offset by 3 because 0, 1, and 2 are reserved indices for “padding,” “start of sequence,” and “unknown.”



• Preparing the data

Listing 4.3 Encoding the integer sequences via multi-hot encoding

```

For training, we need fixed size of given set.
[3, 1, 3, 2, 5]
[3, 1, 1, 1] } we should make these dimension same.
[1, 3, 1, 3, 5, 6]

import numpy as np
def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension)) ← Creates an all-zero matrix
    for i, sequence in enumerate(sequences):          of shape (len(sequences),
        for j in sequence:                           dimension)
            results[i, j] = 1. ← Sets specific indices
    return results                                    of results[i] to 1s

x_train = vectorize_sequences(train_data)           ← Vectorized
x_test = vectorize_sequences(test_data)             training data
                                                    ← Vectorized test data

```

* multi-hot encoding

Vectorize labels

```

y_train = np.asarray(train_labels).astype('float32') ← float int type
y_test = np.asarray(test_labels).astype('float32')

```



Step 1. prepare our data

Step 2. building NN model

• Building your network

Listing 4.4 Model definition

```
from tensorflow import keras
from tensorflow.keras import layers

model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
```

16 output units

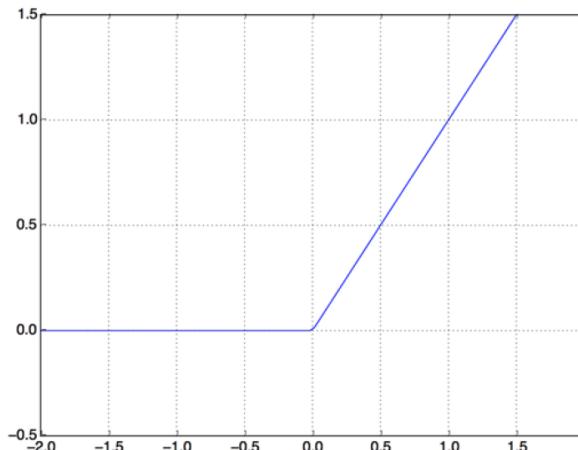
max(0,x)

→ classifying Pos/Neg

(-∞, ∞) → [0, 1]

*Neg ↑
↑ Pos*

Rectified linear unit (relu)



Sigmoid

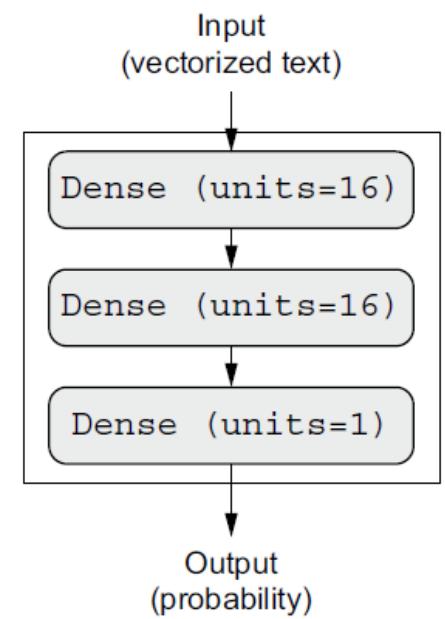
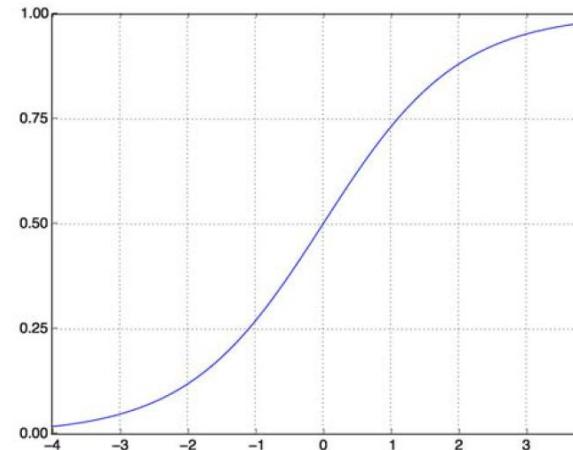
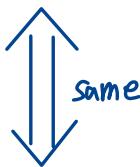


Figure 4.1 The three-layer model



Listing 4.5 Compiling the model

```
model.compile(optimizer="rmsprop",  
              loss="binary_crossentropy",  
              metrics= ["accuracy"] )
```



- **Configuring optimizer, losses and metrics**

```
model.compile(optimizer=optimizers.RMSprop(lr=0.001),  
              loss=losses.binary_crossentropy,  
              metrics=[metrics.binary_accuracy])
```



- **Validating your approach**

Listing 4.6 Setting aside a validation set

```
x_val = x_train[:10000]
partial_x_train = x_train[10000:]
y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```

Listing 4.7 Training your model

```
history = model.fit(partial_x_train,
                     partial_y_train,
                     epochs=20,
                     batch_size=512,
                     validation_data=(x_val, y_val))
```



Listing 4.8 Plotting the training and validation loss

```
import matplotlib.pyplot as plt
history_dict = history.history
loss_values = history_dict["loss"]
val_loss_values = history_dict["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, "bo", label="Training loss") ← "bo" is for
plt.plot(epochs, val_loss_values, "b", label="Validation loss") ← "b" is for
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
```

Listing 4.9 Plotting the training and validation accuracy

```
plt.clf() ← Clears the figure
acc = history_dict["accuracy"]
val_acc = history_dict["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



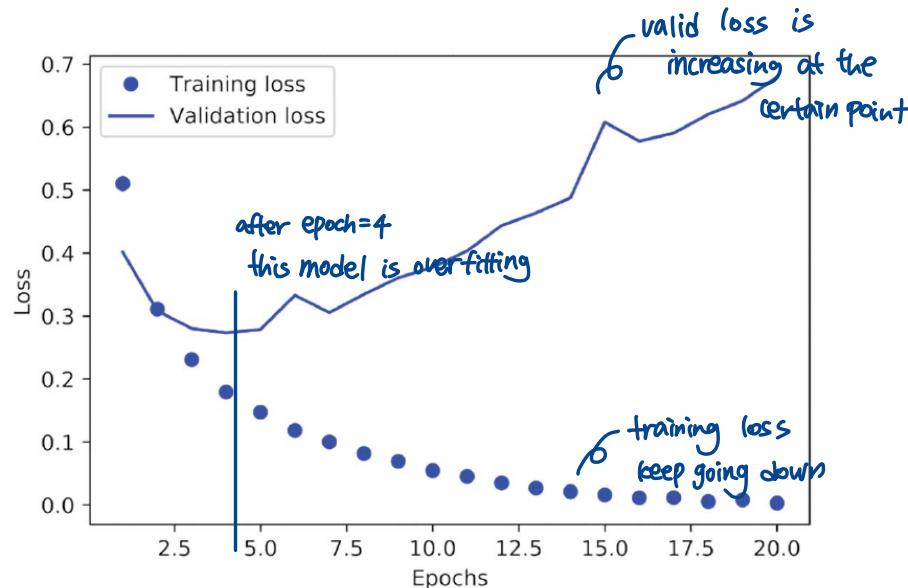


Figure 4.4 Training and validation loss

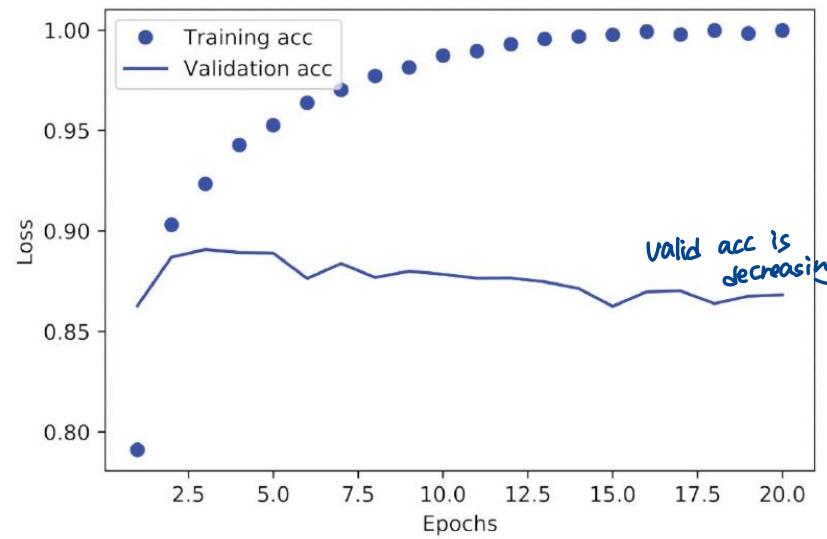


Figure 4.5 Training and validation accuracy

Overfitting: after the fourth epoch, you're overoptimizing on the training data



Listing 4.10 Retraining a model from scratch

```
model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.fit(x_train, y_train, epochs=4, batch_size=512)
results = model.evaluate(x_test, y_test)
```

- **The final results are as follows:**

```
>>> results
[0.2929924130630493, 0.8832799999999995]
```

loss for test set

acc for test

The first number, 0.29, is the test loss, and the second number, 0.88, is the test accuracy.



Using a trained model to generate predictions on new data

```
>>> model.predict(x_test)
array([[ 0.98006207]
       [ 0.99758697]
       [ 0.99975556]
       ...
       [ 0.82167041]
       [ 0.02885115]
       [ 0.65371346]], dtype=float32)
```



Exercises :D

- Try using one or three hidden layers, and see how doing so affects validation and test accuracy.
- Try using layers with more hidden units or fewer hidden units
- Try using the mse loss function instead of binary_crossentropy.
- Try using the tanh activation instead of relu.



Classifying newswires: a multiclass classification example

The Reuters dataset:

- **Dataset for text classification (46 different topics)**
- **8,982 training and 2,246 test examples**



Listing 4.11 Loading the Reuters dataset

```
from tensorflow.keras.datasets import reuters
(train_data, train_labels), (test_data, test_labels) = reuters.load_data(
    num_words=10000)
    using most frequent 10,000 words
```

Listing 4.12 Decoding newswires back to text

```
word_index = reuters.get_word_index()
reverse_word_index = dict([
    (value, key) for (key, value) in word_index.items()])
decoded_newswire = " ".join([
    reverse_word_index.get(i - 3, "?") for i in train_data[0]])
```

Note that the indices are offset by 3 because 0, 1, and 2 are reserved indices for “padding,” “start of sequence,” and “unknown.”



Listing 4.13 Encoding the input data

```
x_train = vectorize_sequences(train_data)      ← Vectorized training data
x_test = vectorize_sequences(test_data)        ← Vectorized test data
```

Listing 4.14 Encoding the labels

```
def to_one_hot(labels, dimension=46):
    results = np.zeros((len(labels), dimension))
    for i, label in enumerate(labels):
        results[i, label] = 1.
    return results
y_train = to_one_hot(train_labels)      ← Vectorized training labels
y_test = to_one_hot(test_labels)        ← Vectorized test labels
```

46

	env	...	sports	...
0	0 0	...	1	...
1	0 0	1	...	0

Can be done using built in function: `to_categorical()`

사용자 정의 함수 안내판

```
from tensorflow.keras.utils import to_categorical
y_train = to_categorical(train_labels)
y_test = to_categorical(test_labels)
```



Building our model

Listing 4.15 Model definition

Listing 4.16 Compiling the model

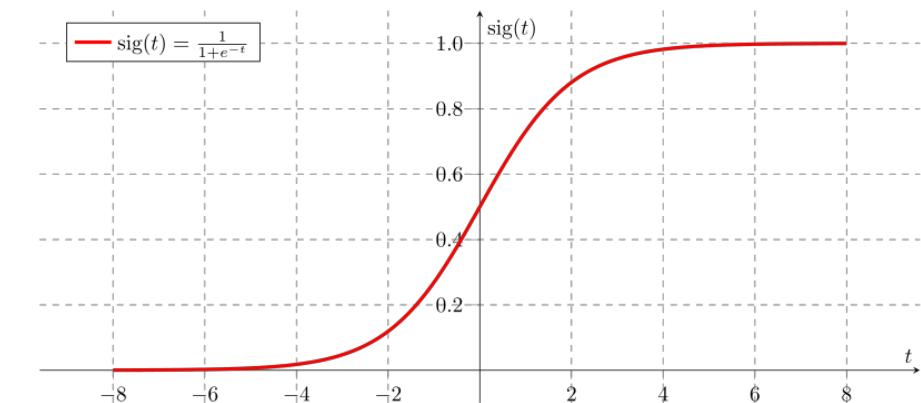
```
model.compile(optimizer="rmsprop",
              loss="categorical_crossentropy",
              metrics= ["accuracy"] )
```



For binary classification, we use sigmoid with single dense output layer

```
model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

$0 \rightarrow \text{Negative}$
 $1 \rightarrow \text{Positive}$



If we use Dense layer with output 2, natural extension should be..

$$\begin{Bmatrix} \frac{e^{-t}}{1 + e^{-t}} \\ \frac{1}{1 + e^{-t}} \end{Bmatrix} = \begin{Bmatrix} \frac{e^{-z_1}}{e^{-z_1} + e^{-z_2}} \\ \frac{e^{-z_2}}{e^{-z_1} + e^{-z_2}} \end{Bmatrix}$$

$$\frac{e^{z_1}}{e^{z_1} + e^{z_2}} = \frac{1}{1 + e^{z_2 - z_1}}$$

```
model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(2, activation='softmax'))
```



Softmax function for multi-class classification

Softmax layer ; extension of sigmoid

$$z_1 \xrightarrow{3} \hat{y}_1 = \frac{e^{z_1}}{\sum e^{z_j}} \quad 0.88$$

$$z_2 \xrightarrow{1} \hat{y}_2 = \frac{e^{z_2}}{\sum e^{z_j}} \quad 0.12$$

$$z_3 \xrightarrow{-3} \hat{y}_3 = \frac{e^{z_3}}{\sum e^{z_j}} \quad \approx 0.0$$

Probability shape: $0 \leq \hat{y}_i \leq 1$

$$\sum \hat{y}_i = 1$$

Listing 4.17 Setting aside a validation set

```
x_val = x_train[:1000]
partial_x_train = x_train[1000:]
y_val = y_train[:1000]
partial_y_train = y_train[1000:]
```

Listing 4.18 Training the model

```
history = model.fit(partial_x_train,
                     partial_y_train,
                     epochs=20,
                     batch_size=512,
                     validation_data=(x_val, y_val))
```



Listing 4.19 Plotting the training and validation loss

```

loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(loss) + 1)
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()

```

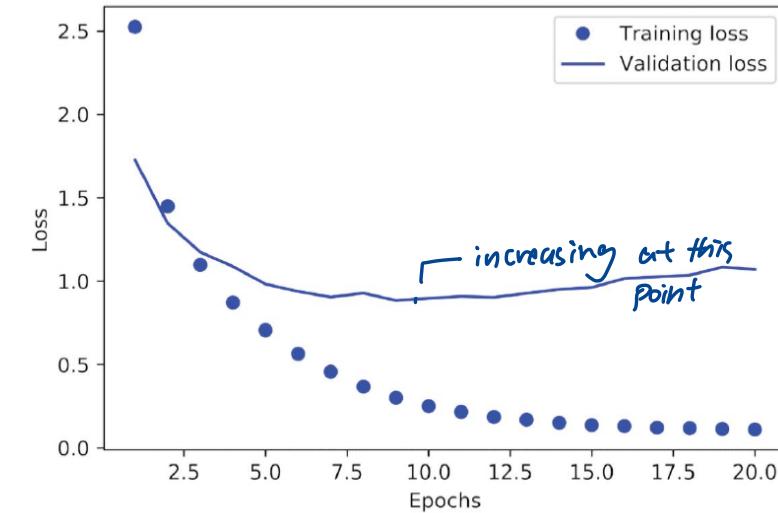


Figure 4.6 Training and validation loss

Listing 4.20 Plotting the training and validation accuracy

```

plt.clf()      ← Clears the figure
acc = history.history["accuracy"]
val_acc = history.history["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training accuracy")
plt.plot(epochs, val_acc, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()

```

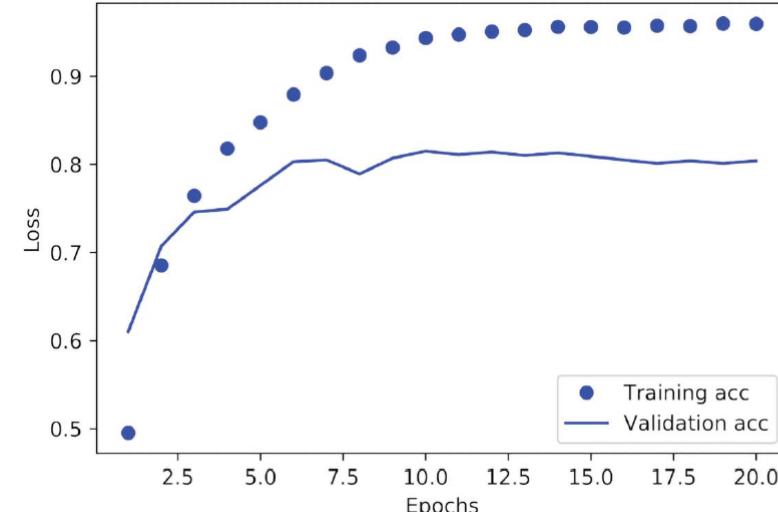


Figure 4.7 Training and validation accuracy

The model begins to overfit after nine epochs.



Listing 4.21 Retraining a model from scratch

```
model = keras.Sequential([
    layers.Dense(64, activation="relu"),
    layers.Dense(64, activation="relu"),
    layers.Dense(46, activation="softmax")
])
model.compile(optimizer="rmsprop",
              loss="categorical_crossentropy",
              metrics=["accuracy"])
model.fit(x_train,
          y_train,
          epochs=9,
          batch_size=512)
results = model.evaluate(x_test, y_test)
```



Final result

```
>>> results           acc: 99.7%  
[0.9565213431445807, 0.79697239536954589]
```

Result from random guess

```
>>> import copy  
>>> test_labels_copy = copy.copy(test_labels)  
>>> np.random.shuffle(test_labels_copy)  
>>> hits_array = np.array(test_labels) == np.array(test_labels_copy)  
>>> hits_array.mean()  
0.18655387355298308
```

A random classifier would score around 19% classification accuracy, so the results of our model seem pretty good in that light.



Generating predictions on new data

```
predictions = model.predict(x_test)
```

Each entry in “predictions” is a vector of length 46:

```
>>> predictions[0].shape  
(46, )
```

The coefficients in this vector sum to 1, as the model use softmax.

```
>>> np.sum(predictions[0])  
1.0
```

The largest entry is the predicted class:

```
>>> np.argmax(predictions[0])  
4
```



A different way to handle the labels and the loss

Another way to encode the labels would be to cast them as an integer tensor, like this:

```
y_train = np.array(train_labels)  
y_test = np.array(test_labels)
```

The only thing this approach would change is the choice of the loss function, categorical_crossentropy -> sparse_categorical_crossentropy

```
model.compile(optimizer="rmsprop",  
              loss="sparse_categorical_crossentropy",  
              metrics=["accuracy"] )
```



The importance of having sufficiently large intermediate layers

Listing 4.22 A model with an information bottleneck

```
model = keras.Sequential([
    layers.Dense(64, activation="relu"),
    layers.Dense(4, activation="relu"), 64→4
    layers.Dense(46, activation="softmax")
])
model.compile(optimizer="rmsprop",
              loss="categorical_crossentropy",
              metrics=["accuracy"])
model.fit(partial_x_train,
          partial_y_train,
          epochs=20,
          batch_size=128,
          validation_data=(x_val, y_val))
```

The model now peaks at ~71% validation accuracy, an 8% absolute drop.



Exercises :D

- **Try using larger or smaller layers**
- **Try using a single hidden layer, or three hidden layers**



Predicting house prices: a regression example

The Boston Housing Price dataset:

- **404 training samples and 102 test samples with 13 features**

Listing 4.23 Loading the Boston housing dataset

```
from tensorflow.keras.datasets import boston_housing
(train_data, train_targets), (test_data, test_targets) = (
    boston_housing.load_data())
```



Listing 4.24 Normalizing the data

```
mean = train_data.mean(axis=0)
train_data -= mean
std = train_data.std(axis=0)
train_data /= std
test_data -= mean
test_data /= std
```

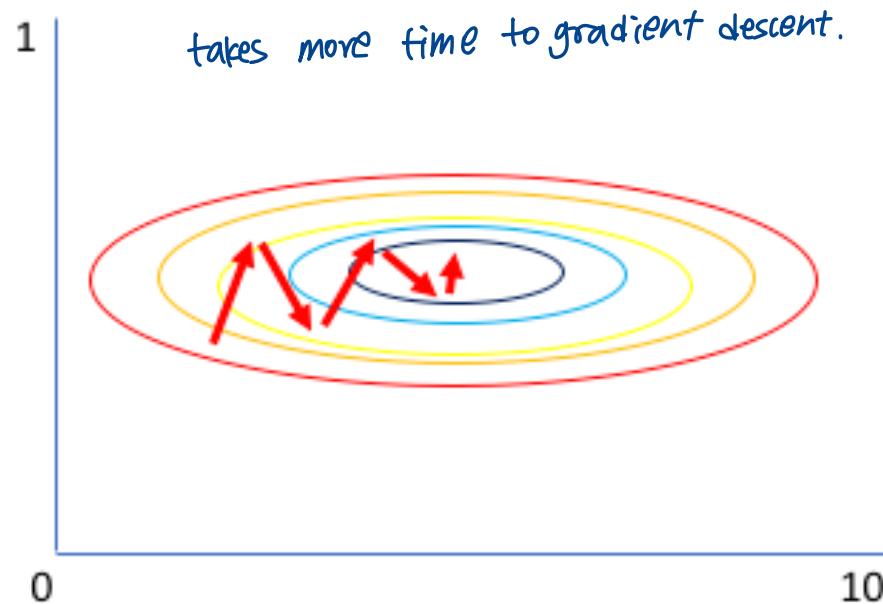
Listing 4.25 Model definition

```
def build_model():
    model = keras.Sequential([
        layers.Dense(64, activation="relu"),
        layers.Dense(64, activation="relu"),
        layers.Dense(1) for regression no need activation
    ])
    model.compile(optimizer="rmsprop", loss="mse", metrics=[ "mae" ] )
    return model
```

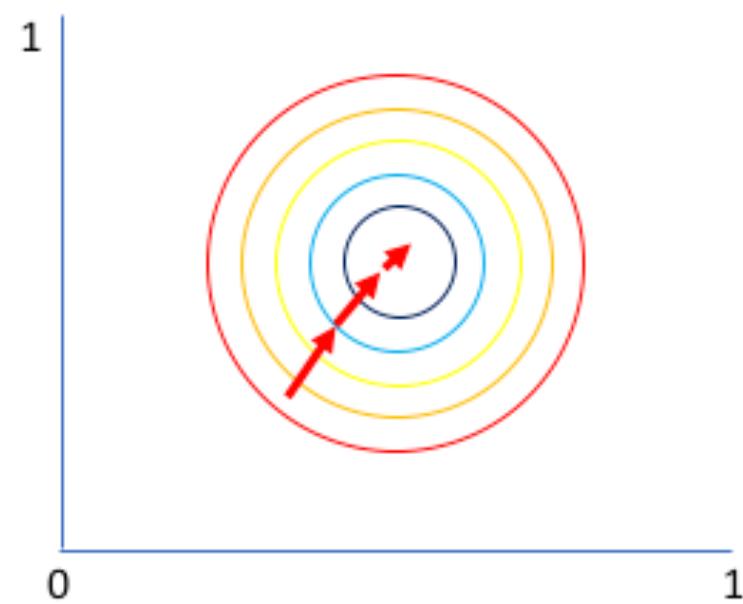
Because we need to instantiate
the same model multiple times,
we use a function to construct it.



Why normalize?



Gradient of larger parameter
dominates the update



Both parameters can be
updated in equal proportions



Validating your approach using K-fold validation

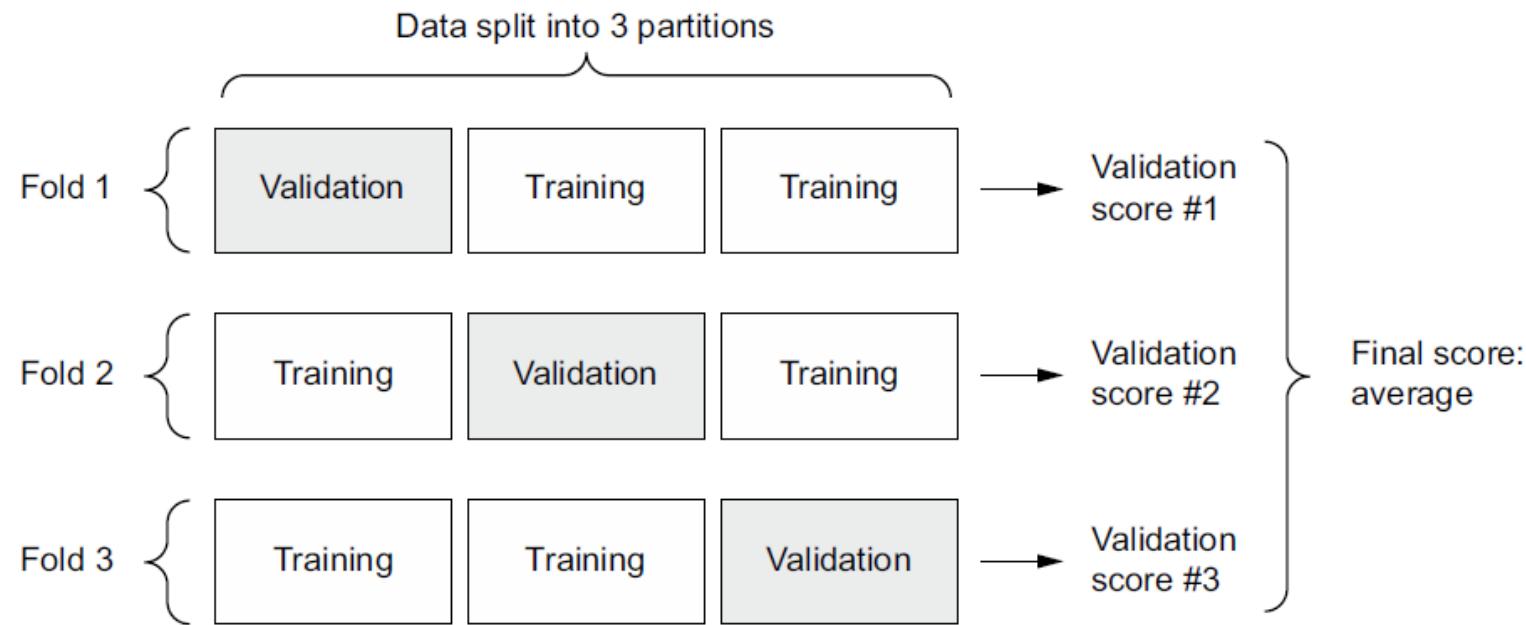


Figure 4.8 K-fold cross-validation with $K=3$



Listing 4.26 K-fold validation

```
k = 4
num_val_samples = len(train_data) // k
num_epochs = 100
all_scores = []
for i in range(k):
    print(f"Processing fold #{i}")
    val_data = train_data[i * num_val_samples: (i + 1) * num_val_samples] ←
    val_targets = train_targets[i * num_val_samples: (i + 1) * num_val_samples]
    partial_train_data = np.concatenate(←
        [train_data[:i * num_val_samples],
         train_data[(i + 1) * num_val_samples:]], ←
        axis=0)
    partial_train_targets = np.concatenate(←
        [train_targets[:i * num_val_samples],
         train_targets[(i + 1) * num_val_samples:]], ←
        axis=0)
    model = build_model()
    model.fit(partial_train_data, partial_train_targets, ←
              epochs=num_epochs, batch_size=16, verbose=0)
    val_mse, val_mae = model.evaluate(val_data, val_targets, verbose=0) ←
    all_scores.append(val_mae)
```

Prepares the validation data: data from partition #k

Prepares the training data: data from all other partitions

Builds the Keras model (already compiled)

Trains the model (in silent mode, verbose = 0)

Evaluates the model on the validation data

```
>>> all_scores
[2.112449, 3.0801501, 2.6483836, 2.4275346]
>>> np.mean(all_scores)
2.5671294
```



Let's try training the model a bit longer: 500 epochs. To keep a record of how well the model does at each epoch, we'll modify the training loop to **save the per epoch validation score log for each fold.**

Listing 4.27 Saving the validation logs at each fold

```

num_epochs = 500
all_mae_histories = []
for i in range(k):
    print(f"Processing fold #{i}")
    val_data = train_data[i * num_val_samples: (i + 1) * num_val_samples]      ← Prepares the validation data: data from partition #k
    val_targets = train_targets[i * num_val_samples: (i + 1) * num_val_samples]
    partial_train_data = np.concatenate(          ← Prepares the training data: data from all other partitions
        [train_data[:i * num_val_samples],
         train_data[(i + 1) * num_val_samples:]],
        axis=0)
    partial_train_targets = np.concatenate(
        [train_targets[:i * num_val_samples],
         train_targets[(i + 1) * num_val_samples:]],
        axis=0)
    model = build_model()          ← Builds the Keras model (already compiled)
    history = model.fit(partial_train_data, partial_train_targets,
                         validation_data=(val_data, val_targets),
                         epochs=num_epochs, batch_size=16, verbose=0)      ← Trains the model (in silent mode, verbose=0)
    mae_history = history.history["val_mae"]      ← list of 4 components
    all_mae_histories.append(mae_history)          ← each is mae_history of 500 epochs.

```

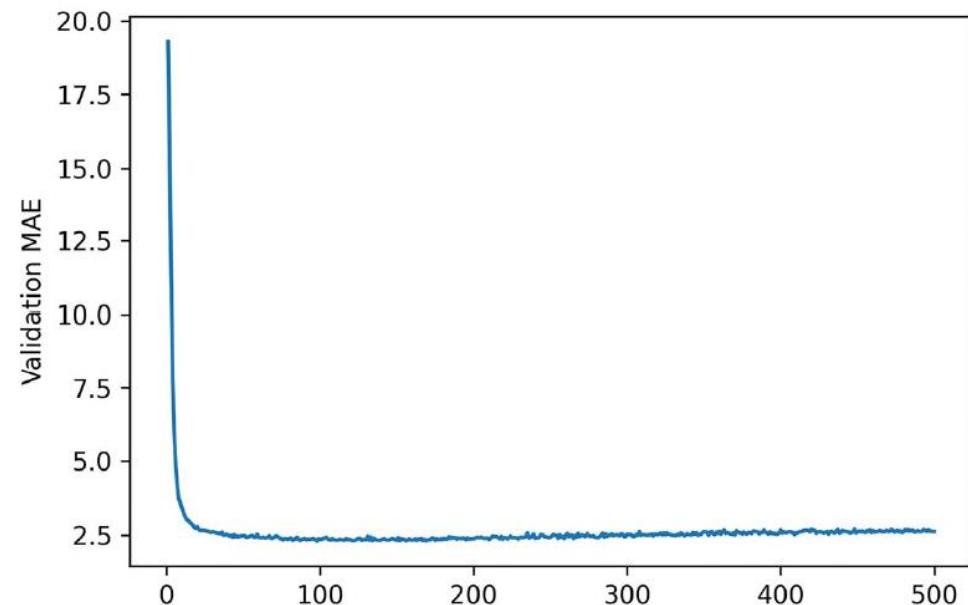


Listing 4.28 Building the history of successive mean K-fold validation scores

```
average_mae_history = [  
    np.mean([x[i] for x in all_mae_histories]) for i in range(num_epochs)]  
    500
```

Listing 4.29 Plotting validation scores

```
plt.plot(range(1, len(average_mae_history) + 1), average_mae_history)  
plt.xlabel("Epochs")  
plt.ylabel("Validation MAE")  
plt.show()
```

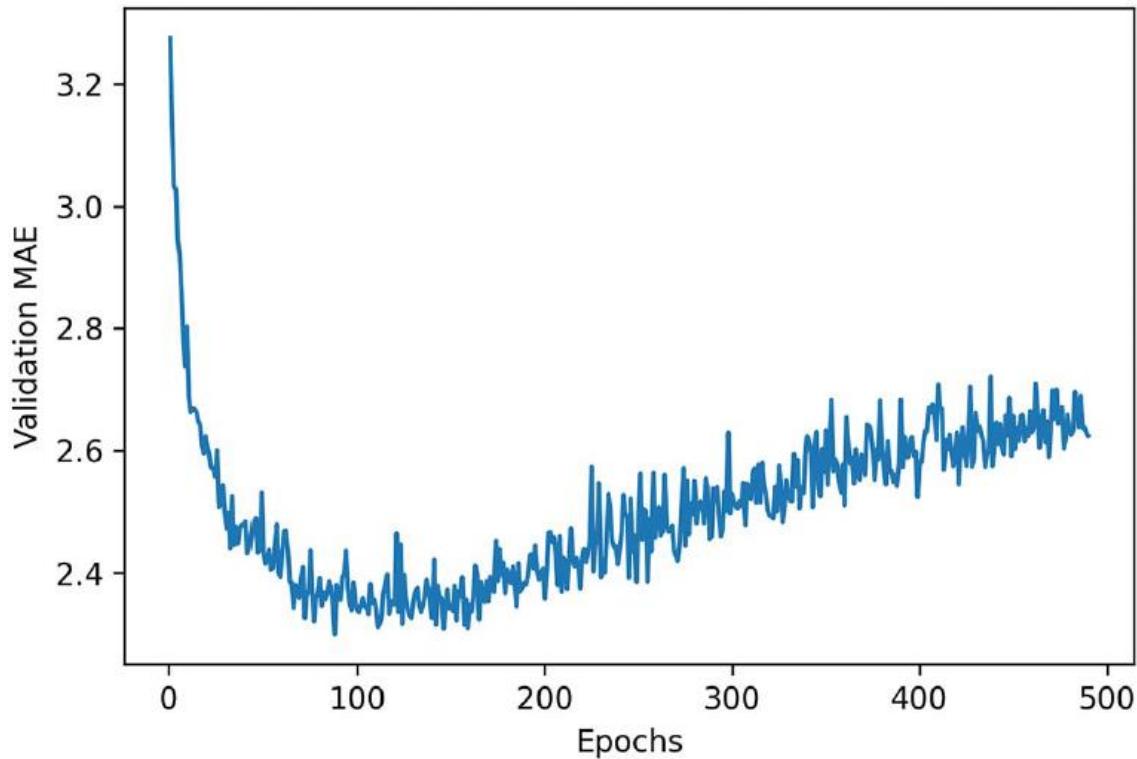


It may be a little difficult to read the plot, due to a scaling issue. Let's omit the first 10 data points.



Listing 4.30 Plotting validation scores, excluding the first 10 data points

```
truncated_mae_history = average_mae_history[10:]
plt.plot(range(1, len(truncated_mae_history) + 1), truncated_mae_history)
plt.xlabel("Epochs")
plt.ylabel("Validation MAE")
plt.show()
```



Validation MAE stops improving significantly after 120–140 epochs

Figure 4.10 Validation MAE by epoch, excluding the first 10 data points



Train a final production model on all of the training data, with the best parameters

Listing 4.31 Training the final model

```
model = build_model()           ← Gets a fresh,  
model.fit(train_data, train_targets,          compiled model  
          epochs=130, batch_size=16, verbose=0)    ← Trains it on the  
test_mse_score, test_mae_score = model.evaluate(test_data, test_targets)  
  
>>> test_mae_score  
2.4642276763916016
```

Generating predictions on new data

```
>>> predictions = model.predict(test_data)  
>>> predictions[0]  
array([9.990133], dtype=float32)
```



chapter04_getting-started-with-neural-networks

File Edit View Insert Runtime Tools Help Cannot save changes

 Share

Table of contents

Classifying newswires: A multiclass classification example

The Reuters dataset

Preparing the data

Building your model

Validating your approach

Generating predictions on new data

A different way to handle the labels and the loss

The importance of having sufficiently large intermediate layers

Further experiments

Wrapping up

Predicting house prices: A regression example

The Boston Housing Price dataset

Preparing the data

Page 11

300

Validating your approach using R fold validation

+ Code + Text Copy to Drive RAM Disk

```
Epoch 13/20
63/63 [=====] - 1s 8ms/step - loss: 0.6067 - accuracy: 0.8454 - val_loss: 1.5636 - val_accuracy: 0.7240
Epoch 14/20
63/63 [=====] - 1s 8ms/step - loss: 0.5687 - accuracy: 0.8520 - val_loss: 1.5497 - val_accuracy: 0.7260
Epoch 15/20
63/63 [=====] - 0s 8ms/step - loss: 0.5406 - accuracy: 0.8594 - val_loss: 1.6150 - val_accuracy: 0.7190
Epoch 16/20
63/63 [=====] - 1s 8ms/step - loss: 0.5133 - accuracy: 0.8690 - val_loss: 1.6685 - val_accuracy: 0.7250
Epoch 17/20
63/63 [=====] - 0s 8ms/step - loss: 0.4922 - accuracy: 0.8752 - val_loss: 1.7528 - val_accuracy: 0.7180
Epoch 18/20
63/63 [=====] - 1s 9ms/step - loss: 0.4682 - accuracy: 0.8820 - val_loss: 1.8296 - val_accuracy: 0.7200
Epoch 19/20
63/63 [=====] - 1s 12ms/step - loss: 0.4489 - accuracy: 0.8849 - val_loss: 1.8731 - val_accuracy: 0.7240
Epoch 20/20
63/63 [=====] - 1s 8ms/step - loss: 0.4308 - accuracy: 0.8898 - val_loss: 1.8995 - val_accuracy: 0.7290
<keras.callbacks.History at 0x7f3d785822d0>
```

Further experiments

Wrapping up

▼ Predicting house prices: A regression example

▼ The Boston Housing Price dataset

Exercises :D

- Try varying the number of layers in the model
- Try varying the number of units per layer



Thank you!

