

CH7. Working with Keras: A deep dive

Il-Youp Kwak
Chung-Ang University



Working with Keras: A deep dive

A spectrum of workflows

Different ways to build Keras models

Using built-in training and evaluation loops

Writing your own training and evaluation loops



A spectrum of workflows

The design of the Keras API: make it easy to get started, yet make it possible to handle high-complexity use cases, only requiring incremental learning at each step

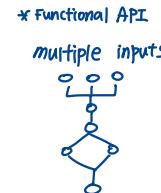
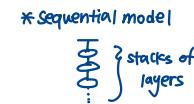
Simple use cases should be easy and approachable, and arbitrarily advanced workflows should be possible



Different ways to build Keras models

Three APIs for building models:

Sequential model: simple stacks of layers



Functional API: focuses on graph-like model architectures

Model subclassing: a low-level option where you write everything yourself.

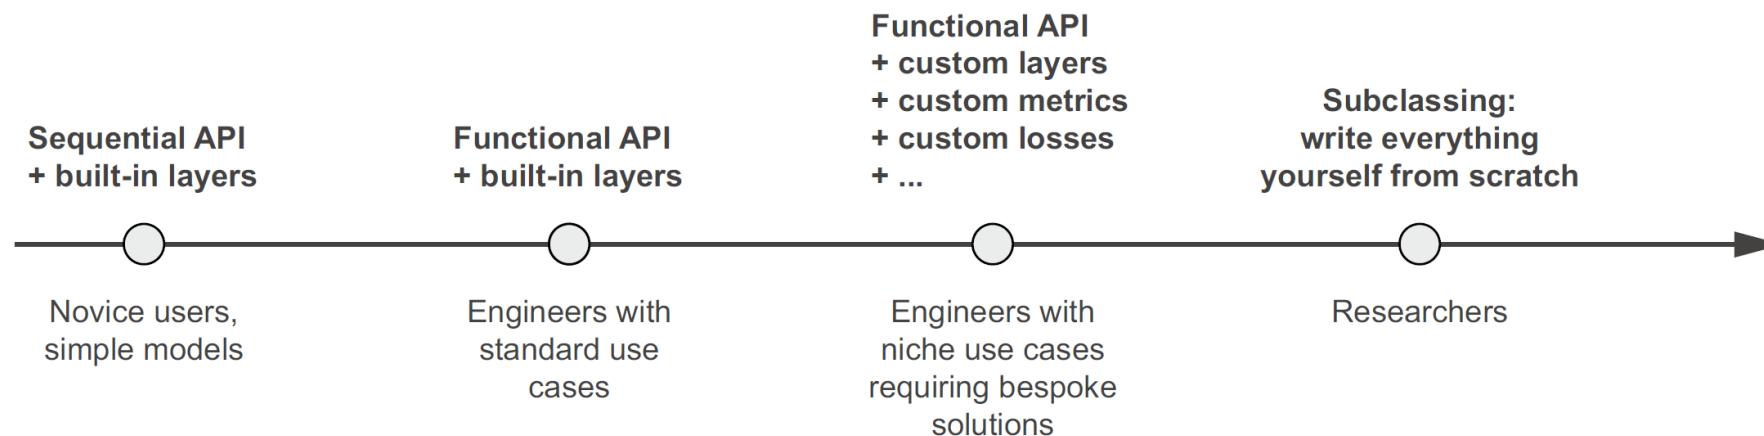


Figure 7.1 Progressive disclosure of complexity for model building

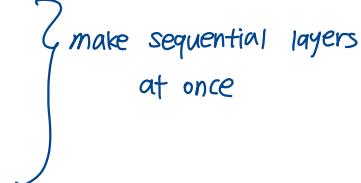


The Sequential model

Listing 7.1 The Sequential class

```
from tensorflow import keras
from tensorflow.keras import layers

model = keras.Sequential([
    layers.Dense(64, activation="relu"),
    layers.Dense(10, activation="softmax")
])
```

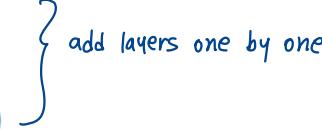


make sequential layers
at once

Build the same model incrementally via the `add()` method

Listing 7.2 Incrementally building a Sequential model

```
model = keras.Sequential()
model.add(layers.Dense(64, activation="relu"))
model.add(layers.Dense(10, activation="softmax"))
```



add layers one by one



Sequential model does not have any weights until you actually call it on some data, or call its build() method with an input shape

If you do not specify input dimension, input weight
then actually cannot calculate weight size.

Listing 7.3 Models that aren't yet built have no weights

```
>>> model.weights
ValueError: Weights for model sequential_1 have not yet been created.
```

At that point, the model isn't built yet.

Listing 7.4 Calling a model for the first time to build it

↳ build(input_shape)가 실행해주면 weights를 할 수 있음

```
>>> model.build(input_shape=(None, 3))
>>> model.weights
[<tf.Variable "dense_2/kernel:0" shape=(3, 64) dtype=float32, ... >,
 <tf.Variable "dense_2/bias:0" shape=(64,) dtype=float32, ... >,
 <tf.Variable "dense_3/kernel:0" shape=(64, 10) dtype=float32, ... >,
 <tf.Variable "dense_3/bias:0" shape=(10,) dtype=float32, ... >]
```

Now you can retrieve the model's weights.

Builds the model—now the model will expect samples of shape (3,). The None in the input shape signals that the batch size could be anything.

$$\begin{aligned} y &= f_2(f_1(x)) \\ f_1(x) &= \sigma(w_1x + b_1) \\ w_1 &\in \mathbb{R}^{3 \times 1}, b_1 \in \mathbb{R}^{1 \times 1} \end{aligned}$$

If don't know about input shape x ,
→ can't define the shape of w_1 .



After the model is built, you can display its contents via the summary()

Listing 7.5 The summary() method

```
>>> model.summary()  
Model: "sequential_1"  


| Layer (type)    | Output Shape | Param #      |                                                 |
|-----------------|--------------|--------------|-------------------------------------------------|
| dense_2 (Dense) | (None, 64)   | 256<br>64x4  | $\sigma(x \cdot w_1 + b_1) = x'$<br>3x64 1x64   |
| dense_3 (Dense) | (None, 10)   | 650<br>65x10 | $\sigma(x' \cdot w_2 + b_2)$<br>1x64 64x10 1x10 |

  
Total params: 906  
Trainable params: 906  
Non-trainable params: 0
```



This model happens to be named “sequential_1.” You can give names to everything in Keras—every model, every layer

Listing 7.6 Naming models and layers with the name argument

```
>>> model = keras.Sequential(name="my_example_model")
>>> model.add(layers.Dense(64, activation="relu", name="my_first_layer"))
>>> model.add(layers.Dense(10, activation="softmax", name="my_last_layer"))
>>> model.build((None, 3))
>>> model.summary()
Model: "my_example_model"
```

Layer (type)	Output Shape	Param #
=====		
my_first_layer (Dense)	(None, 64)	256
=====		
my_last_layer (Dense)	(None, 10)	650
=====		
Total params: 906		
Trainable params: 906		
Non-trainable params: 0		



If you'd like to print summary(), declare the shape of the model's inputs in advance

Listing 7.7 Specifying the input shape of your model in advance

```
model = keras.Sequential()  
model.add(keras.Input(shape=(3,)))  
model.add(layers.Dense(64, activation="relu"))  
  
>>> model.summary()  
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 64)	256

Total params: 256
Trainable params: 256
Non-trainable params: 0

```
>>> model.add(layers.Dense(10, activation="softmax"))  
>>> model.summary()  
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 64)	256
dense_5 (Dense)	(None, 10)	650

Total params: 906
Trainable params: 906
Non-trainable params: 0

Use Input to declare the shape of the inputs. Note that the shape argument must be the shape of each sample, not the shape of one batch.



colab.research.google.com/github/fchollet/deep-learning-with-pyth
news study data science 학회사이트 업무관련 기타 Bookmark M

Pause 00:00:00 Select Area Audio Record Pointer

keras.ipynb 번역 SR Translate scm Data 학술지관련

Share

CO chapter07_working-with-keras.ipynb

File Edit View Insert Runtime Tools Help

Table of contents

+ Code + Text Copy to Drive Connect Editing

Working with Keras: A deep dive

A spectrum of workflows

Different ways to build Keras models

The Sequential model

The Functional API

A simple example

Multi-input, multi-output models

Training a multi-input, multi-output model

The power of the Functional API: Access to layer connectivity

Subclassing the Model class

Rewriting our previous example as a subclassed model

Beware: What subclassed models don't support

Mixing and matching different components

Remember: Use the right tool for the job

This is a companion notebook for the book [Deep Learning with Python, Second Edition](#). For readability, it only contains runnable code blocks and section titles, and omits everything else in the book: text paragraphs, figures, and pseudocode.

If you want to be able to follow what's going on, I recommend reading the notebook side by side with your copy of the book.

This notebook was generated for TensorFlow 2.6.

Working with Keras: A deep dive

A spectrum of workflows

Different ways to build Keras models

The Sequential model

The Sequential class

```
[ ] from tensorflow import keras
from tensorflow.keras import layers
```

colab.research.google.com의 응답을 기다리는 중...

검색하려면 여기에 입력하십시오.

7°C 맑음

오늘 7:56
2022-04-02

The Functional API

Sequential model only express models with **single input and single output**

& model is sequentially defined.

Functional API provide more power to define general model, *graph-like structure*

Listing 7.8 A simple Functional model with two Dense layers

```
inputs = keras.Input(shape=(3,), name="my_input")
features = layers.Dense(64, activation="relu")(inputs)
outputs = layers.Dense(10, activation="softmax")(features)
model = keras.Model(inputs=inputs, outputs=outputs)
```

inputs object holds information about the shape and dtype of the data

```
>>> inputs.shape
(None, 3)
>>> inputs.dtype
float32
```

The model will process batches where each sample has shape (3,). The number of samples per batch is variable (indicated by the None batch size).

These batches will have dtype float32.

symbolic information

Data is not provided.

We call such an object a **symbolic tensor**. It doesn't contain any actual data, but it encodes the specifications of the actual tensors of data that the model will see when you use it.



```
features = layers.Dense(64, activation="relu")(inputs)
```

real tensors of data
or
these symbolic tensors

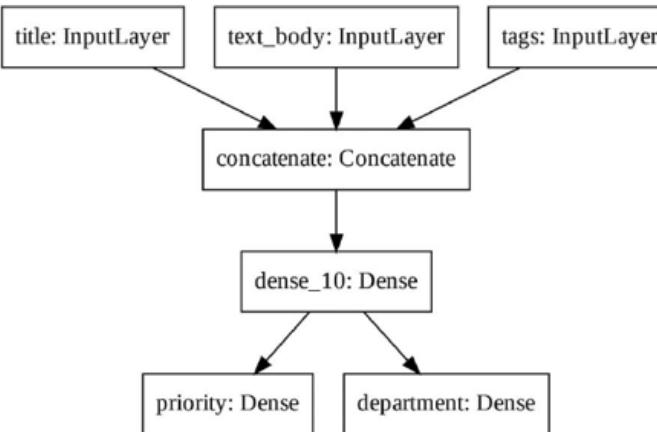
All Keras layers can be called both on real tensors of data and on these symbolic tensors. In the latter case, they return a new symbolic tensor, with updated shape and dtype information:

```
>>> features.shape  
(None, 64)
```

After obtaining the final outputs, we instantiated the model by specifying its inputs and outputs in the Model constructor:
예를 들어 설명하다

```
outputs = layers.Dense(10, activation="softmax")(features)  
model = keras.Model(inputs=inputs, outputs=outputs)
```





Multi-input, Multi-output models

→ cannot define sequential API
→ need to define functional API

Most deep learning models look like graphs.

Consider multi-input, multi-output model:

Listing 7.9 A multi-input, multi-output Functional model

```

vocabulary_size = 10000
num_tags = 100
num_departments = 4

Define model inputs. title = keras.Input(shape=(vocabulary_size,), name="title")
text_body = keras.Input(shape=(vocabulary_size,), name="text_body")
tags = keras.Input(shape=(num_tags,), name="tags")

Combine input features into a single tensor, features, by concatenating them.

features = layers.concatenate([title, text_body, tags]) ← result shape → 20,100 vector
features = layers.Dense(64, activation="relu")(features)

Define model outputs. priority = layers.Dense(1, activation="sigmoid", name="priority")(features)
department = layers.Dense(
    num_departments, activation="softmax", name="department")(features)

model = keras.Model(inputs=[title, text_body, tags], outputs=[priority, department])
  
```

Create the model by specifying its inputs and outputs.

we have 3 inputs & 2 outputs

Apply an intermediate layer to recombine input features into richer representations.



Training a Multi-input, Multi-output model

Listing 7.10 Training a model by providing lists of input and target arrays

```
import numpy as np

num_samples = 1280

Dummy input data | title_data = np.random.randint(0, 2, size=(num_samples, vocabulary_size))
                     text_body_data = np.random.randint(0, 2, size=(num_samples, vocabulary_size))
                     tags_data = np.random.randint(0, 2, size=(num_samples, num_tags))

Dummy target data | priority_data = np.random.random(size=(num_samples, 1))
                     department_data = np.random.randint(0, 2, size=(num_samples, num_departments))

model.compile(optimizer="rmsprop",
              loss=["mean_squared_error", "categorical_crossentropy"],
              metrics=[[ "mean_absolute_error"], [ "accuracy"]])
model.fit([title_data, text_body_data, tags_data],
          [priority_data, department_data],
          epochs=1)
model.evaluate([title_data, text_body_data, tags_data],
               [priority_data, department_data])
priority_preds, department_preds = model.predict(
    [title_data, text_body_data, tags_data])
```

should be new data

dummy data



You can also leverage the names you gave to the Input objects and the output layers, and pass data via dictionaries.

Listing 7.11 Training a model by providing dicts of input and target arrays

```
model.compile(optimizer="rmsprop",
              loss={"priority": "mean_squared_error", "department":
                    "categorical_crossentropy"},
              metrics={"priority": ["mean_absolute_error"], "department":
                    ["accuracy"]})
model.fit({"title": title_data, "text_body": text_body_data,
           "tags": tags_data},
           {"priority": priority_data, "department": department_data},
           epochs=1)
model.evaluate({"title": title_data, "text_body": text_body_data,
                "tags": tags_data},
                {"priority": priority_data, "department": department_data})
priority_preds, department_preds = model.predict(
    {"title": title_data, "text_body": text_body_data, "tags": tags_data})
```



The Power of the functional API: Access to layer connectivity

A Functional model is an explicit graph data structure. This makes it possible to inspect how layers are connected and reuse previous graph nodes as part of new models.

After carefully inspecting the model,
you can actually reuse previous graph nodes as part of new models.

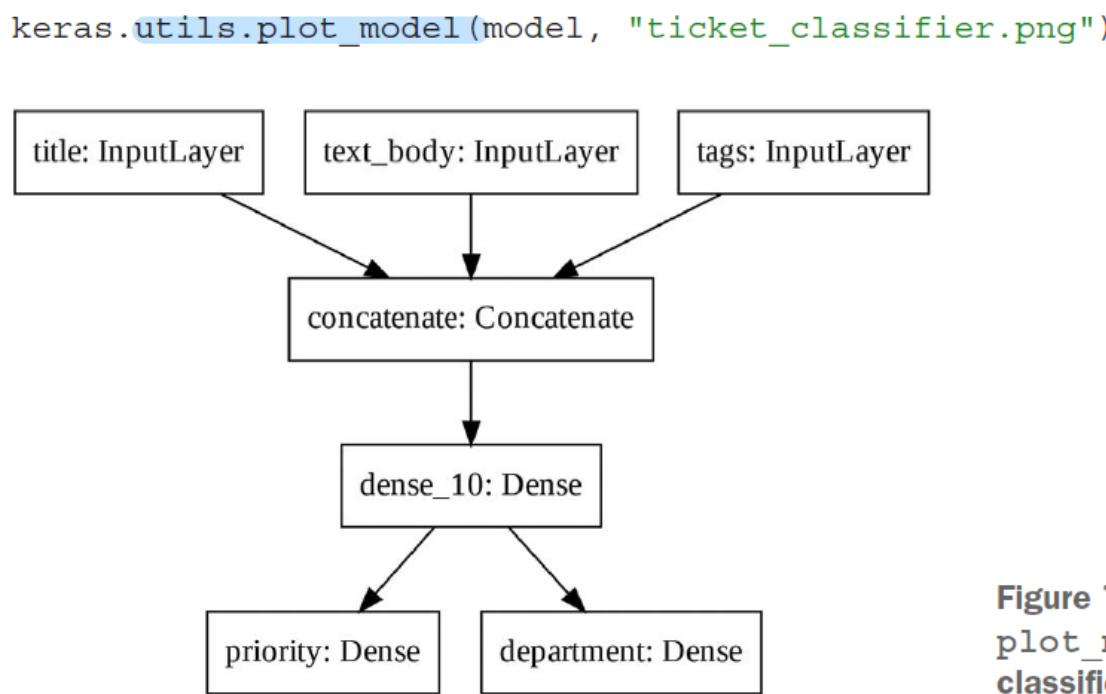


Figure 7.2 Plot generated by `plot_model()` on our ticket classifier model

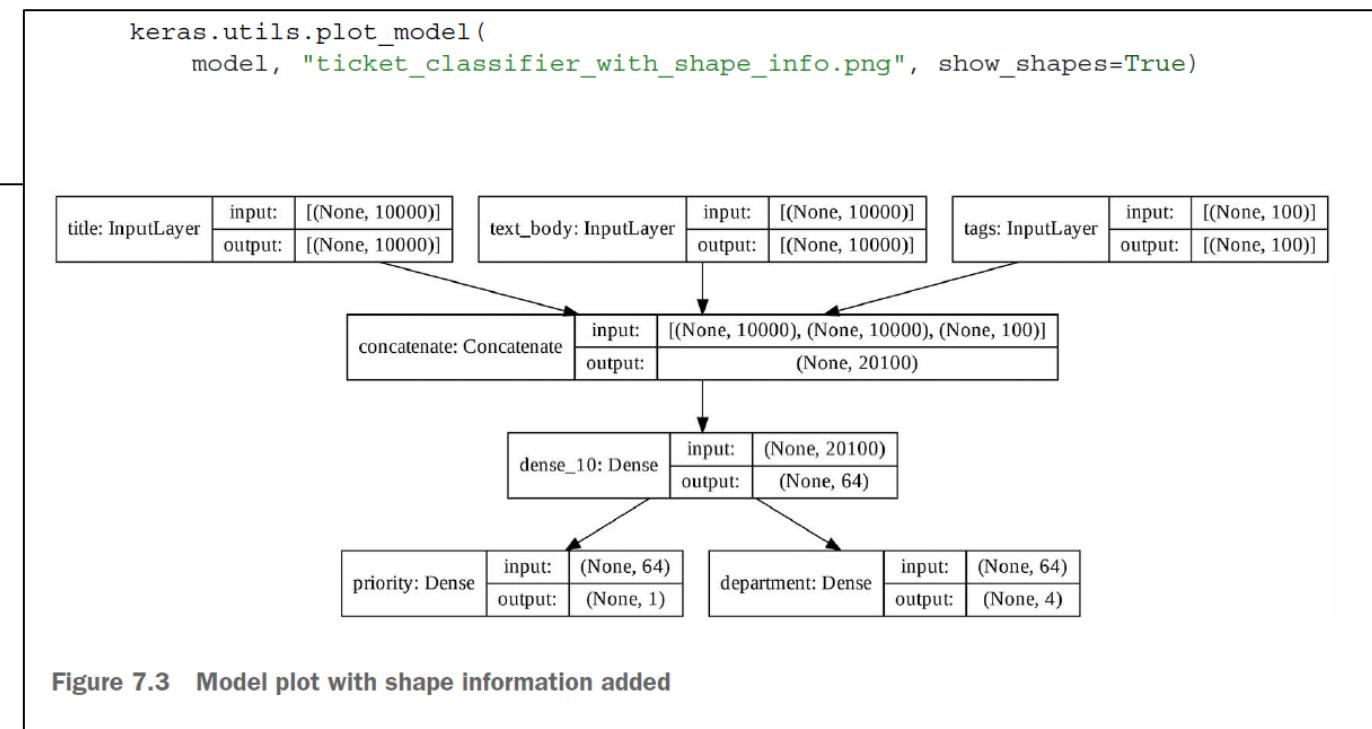


Figure 7.3 Model plot with shape information added



Access to layer connectivity also means that you can inspect and reuse individual nodes (layer calls) in the graph.

Listing 7.12 Retrieving the inputs or outputs of a layer in a Functional model

```
>>> check layers
>>> model.layers
[0]<tensorflow.python.keras.engine.input_layer.InputLayer at 0x7fa963f9d358>,
[1]<tensorflow.python.keras.engine.input_layer.InputLayer at 0x7fa963f9d2e8>,
[2]<tensorflow.python.keras.engine.input_layer.InputLayer at 0x7fa963f9d470>,
[3]<tensorflow.python.keras.layers.merge.Concatenate at 0x7fa963f9d860>,
  <tensorflow.python.keras.layers.core.Dense at 0x7fa964074390>,
  <tensorflow.python.keras.layers.core.Dense at 0x7fa963f9d898>,
  <tensorflow.python.keras.layers.core.Dense at 0x7fa963f95470>
>>> model.layers[3]concatenate layer.input
[<tf.Tensor "title:0" shape=(None, 10000) dtype=float32>,
 <tf.Tensor "text_body:0" shape=(None, 10000) dtype=float32>,
 <tf.Tensor "tags:0" shape=(None, 100) dtype=float32>]
>>> model.layers[3].output
<tf.Tensor "concatenate(concat:0" shape=(None, 20100) dtype=float32>
```

Enables us to do feature extraction, creating models that reuse intermediate features from another model.

Listing 7.13 Creating a new model by reusing intermediate layer outputs

```
features = model.layers[4].output
difficult = layers.Dense(3, activation="softmax", name="difficulty") (features)#new dense layer

new_model = keras.Model(
    inputs=[title, text_body, tags],
    outputs=[priority, department, difficulty])output is added
```

layers[4] is our intermediate Dense layer

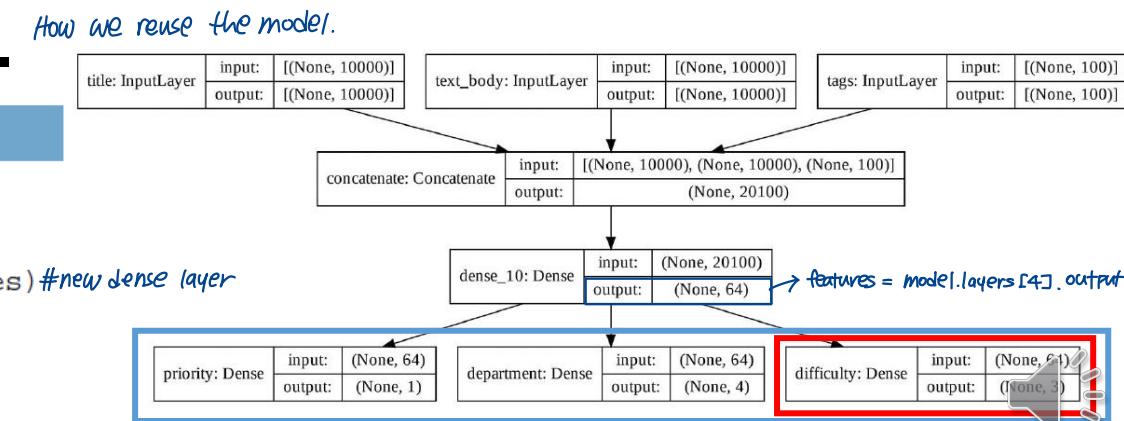


Figure 7.4 Plot of our new model

Subclassing the Model class

In the `__init__()` method, define the layers the model will use

In the `call()` method, define the forward pass of the model

Instantiate your subclass, and call it on data to create its weights

Listing 7.14 A simple subclassed model

```
class CustomerTicketModel(keras.Model):
    → class created by keras
    → right; constructor
    def __init__(self, num_departments):
        super().__init__()
        self.concat_layer = layers.concatenate()
        self.mixing_layer = layers.Dense(64, activation="relu")
        self.priority_scorer = layers.Dense(1, activation="sigmoid")
        self.department_classifier = layers.Dense(
            num_departments, activation="softmax")
    → should be dictionary type
    def call(self, inputs):
        title = inputs["title"]
        text_body = inputs["text_body"]
        tags = inputs["tags"]

        features = self.concat_layer([title, text_body, tags])
        features = self.mixing_layer(features)
        priority = self.priority_scorer(features)
        department = self.department_classifier(features)
        return priority, department
```

Don't forget to
call the `super()`
constructor!

Define
sublayers
in the
constructor.

Define the forward
pass in the `call()`
method.



Once you've defined the model, you can instantiate it.

```
model = CustomerTicketModel(num_departments=4)

priority, department = model(
    {"title": title_data, "text_body": text_body_data, "tags": tags_data})
```

You can compile and train a Model subclass just like a Sequential or Functional model:

```
model.compile(optimizer="rmsprop",
              loss=["mean_squared_error", "categorical_crossentropy"],
              metrics=[["mean_absolute_error"], ["accuracy"]])

model.fit({"title": title_data,
           "text_body": text_body_data,
           "tags": tags_data},
           [priority_data, department_data], ←
           epochs=1)

model.evaluate({"title": title_data,
                "text_body": text_body_data,
                "tags": tags_data},
                [priority_data, department_data])

priority_preds, department_preds = model.predict({"title": title_data,
                                                 "text_body": text_body_data,
                                                 "tags": tags_data})
```

The structure of the input data must match exactly what is expected by the call() method—here, a dict with keys title, text_body, and tags.

The structure of what you pass as the loss and metrics arguments must match exactly what gets returned by call()—here, a list of two elements.

The structure of the target data must match exactly what is returned by the call() method—here, a list of two elements.



Beware: what subclassed models don't support

This freedom comes at a cost: with subclassed models, you are responsible for more of the model logic, which means your potential error surface is much larger



colab.research.google.com/github/fchollet/deep-learning-with-python-notebooks/blob/master/chapter07_working-with-keras.ipynb#scrollTo=SiE7sDfb6ybA

news study data science 학회사이트 업무관련 기타 Bookmark Manager Gmail YouTube 지도 뉴스 번역 SR Translate scm Data 학술지관련 » 기타

Share

Table of contents

Working with Keras: A deep dive

- A spectrum of workflows
- Different ways to build Keras models
 - The Sequential model
 - The Functional API
 - A simple example
 - Multi-input, multi-output models
 - Training a multi-input, multi-output model
 - The power of the Functional API: Access to layer connectivity
- Subclassing the Model class
 - Rewriting our previous example as a subclassed model
 - Beware: What subclassed models don't support
 - Mixing and matching different components
 - Remember: Use the right tool for the job

keras.utils.plot_model(new_model, "updated_ticket_classifier.png", show_shapes=True)

```
graph TD; title[title] --- conc[concatenate_1]; text_body[text_body] --- conc; tags[tags] --- conc; conc --> dense[dense_11]; dense --> priority[Dense priority]; dense --> department[Dense department]; dense --> difficulty[Dense difficulty]
```

RAM Disk ✓ Editing

Subclassing the Model class

Rewriting our previous example as a subclassed model

A simple subclassed model

```
class CustomerTicketModel(keras.Model):
```

검색하려면 여기에 입력하십시오.

9:00 6°C 맑음

오늘 8:40
2022-04-02

Mixing and matching different components

All models in the Keras API can smoothly interoperate with each other, whether they're Sequential models, Functional models, or subclassed models

Listing 7.15 Creating a Functional model that includes a subclassed model

```
class Classifier(keras.Model):
    ↪ Subclass model

    def __init__(self, num_classes=2):
        super().__init__()
        if num_classes == 2:
            num_units = 1
            activation = "sigmoid"
        else:
            num_units = num_classes
            activation = "softmax"
        self.dense = layers.Dense(num_units, activation=activation)

    def call(self, inputs):
        return self.dense(inputs)

inputs = keras.Input(shape=(3,))
features = layers.Dense(64, activation="relu")(inputs)
outputs = Classifier(num_classes=10)(features)
model = keras.Model(inputs=inputs, outputs=outputs)
```



Inversely, you can use a Functional model as part of a subclassed layer or model

Listing 7.16 Creating a subclassed model that includes a Functional model

```
inputs = keras.Input(shape=(64,))
outputs = layers.Dense(1, activation="sigmoid")(inputs)
binary_classifier = keras.Model(inputs=inputs, outputs=outputs)

class MyModel(keras.Model):

    def __init__(self, num_classes=2):
        super().__init__()
        self.dense = layers.Dense(64, activation="relu")
        self.classifier = binary_classifier ← functional model in constructor

    def call(self, inputs):
        features = self.dense(inputs)
        return self.classifier(features)

model = MyModel()
```



Using built-in training and evaluation loops

Listing 7.17 The standard workflow: `compile()`, `fit()`, `evaluate()`, `predict()`

```

from tensorflow.keras.datasets import mnist
create MNIST model
using functional API
def get_mnist_model():
    inputs = keras.Input(shape=(28 * 28,))
    features = layers.Dense(512, activation="relu")(inputs)
    features = layers.Dropout(0.5)(features)
    outputs = layers.Dense(10, activation="softmax")(features)
    model = keras.Model(inputs, outputs)
    return model
Prepare our
training and test
(data sets) also (labels)
(images, labels), (test_images, test_labels) = mnist.load_data()
images = images.reshape((60000, 28 * 28)).astype("float32") / 255
divide for
normalization
test_images = test_images.reshape((10000, 28 * 28)).astype("float32") / 255
train_images, val_images = images[10000:], images[:10000]
train_labels, val_labels = labels[10000:], labels[:10000]
model = get_mnist_model()
model.compile(optimizer="rmsprop",
              loss="sparse_categorical_crossentropy",
              metrics=["accuracy"])
built-in
training loop
model.fit(train_images, train_labels,
          epochs=3,
          validation_data=(val_images, val_labels))
built-in
evaluation loop
test_metrics = model.evaluate(test_images, test_labels)
predictions = model.predict(test_images)
Use evaluate() to
compute the loss and
metrics on new data.
Use predict() to compute
classification probabilities
on new data.
Create a model (we factor this
into a separate function so as
to reuse it later).
Load your data, reserving
some for validation.
Compile the model by
specifying its optimizer, the
loss function to minimize,
and the metrics to monitor.
Use fit() to train the
model, optionally
providing validation data
to monitor performance
on unseen data.

```

Provide a custom metric

after defining own versions

Pass **callbacks** to the `fit()` to schedule actions at specific points during training. cf) early stopping



Writing your own metrics

A Keras metric is a subclass of the `keras.metrics.Metric` class

Write the state-update logic, which happens in the `update_state()` method.

Listing 7.18 Implementing a custom metric by subclassing the Metric class

```
import tensorflow as tf

class RootMeanSquaredError(keras.metrics.Metric):
    def __init__(self, name="rmse", **kwargs):
        super().__init__(name=name, **kwargs)
        self.mse_sum = self.add_weight(name="mse_sum", initializer="zeros")
        self.total_samples = self.add_weight(
            name="total_samples", initializer="zeros", dtype="int32")

    def update_state(self, y_true, y_pred, sample_weight=None):
        calculate MSE
        y_true = tf.one_hot(y_true, depth=tf.shape(y_pred)[1])
        mse = tf.reduce_sum(tf.square(y_true - y_pred))
        self.mse_sum.assign_add(mse) → It will add MSE one by one at every training stage
        num_samples = tf.shape(y_pred)[0] → update
        self.total_samples.assign_add(num_samples) → Add # samples one by one, if batch size → 32 we add 32 times for total samples

    def result(self):
        return tf.sqrt(self.mse_sum / tf.cast(self.total_samples, tf.float32))

    def reset_state(self):
        self.mse_sum.assign(0.)
        self.total_samples.assign(0.)
```

Define the state variables in the constructor. Like for layers, you have access to the `add_weight()` method.

To match our MNIST model, we expect categorical predictions and integer labels.

Subclass the Metric class.

epoch 1
→→→→→→→→ ...
each update

Implement the state update logic in `update_state()`. The `y_true` argument is the targets (or labels) for one batch, while `y_pred` represents the corresponding predictions from the model. You can ignore the `sample_weight` argument—we won't use it here.

return current metric:

```
def result(self):
    return tf.sqrt(self.mse_sum / tf.cast(self.total_samples, tf.float32)) → For each epoch, the current MSE will be updated
    by this method.

def reset_state(self):
    self.mse_sum.assign(0.) → When I start new epoch, should reset the state.
    self.total_samples.assign(0.)
```



Custom metrics can be used just like built-in ones

```
model = get_mnist_model()
model.compile(optimizer="rmsprop",
              loss="sparse_categorical_crossentropy",
              metrics=["accuracy", RootMeanSquaredError()])
model.fit(train_images, train_labels,
          epochs=3,
          validation_data=(val_images, val_labels))
test_metrics = model.evaluate(test_images, test_labels)
```



Table of contents

- Working with Keras: A deep dive
- A spectrum of workflows
- Different ways to build Keras models
 - The Sequential model
 - The Functional API
 - A simple example
 - Multi-input, multi-output models
 - Training a multi-input, multi-output model
 - The power of the Functional API: Access to layer connectivity
- Subclassing the Model class
 - Rewriting our previous example as a subclassed model
 - Beware: What subclassed models don't support**
- Mixing and matching different components
- Remember: Use the right tool for the job

+ Code + Text Copy to Drive RAM Disk

```
    "tags": tags_data})  
40/40 [=====] - 1s 8ms/step - loss: 35.1242 - output_1_loss: 0.3360 - output_2_loss: 34.7882 - output_1_mean_absolute_error:  
40/40 [=====] - 0s 6ms/step - loss: 35.1771 - output_1_loss: 0.3426 - output_2_loss: 34.8345 - output_1_mean_absolute_error:  
  
Beware: What subclassed models don't support
```

▼ Mixing and matching different components

Creating a Functional model that includes a subclassed model

```
[ ] class Classifier(keras.Model):  
    def __init__(self, num_classes=2):  
        super().__init__()  
        if num_classes == 2:  
            num_units = 1  
            activation = "sigmoid"  
        else:  
            num_units = num_classes  
            activation = "softmax"  
        self.dense = layers.Dense(num_units, activation=activation)  
  
    def call(self, inputs):  
        return self.dense(inputs)
```

3s completed at 8:43 PM

Using callbacks

in training stage

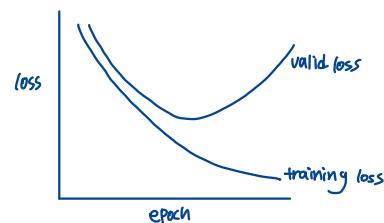
Launching a training run on a large dataset for tens of epochs using `model.fit()` can be a bit like launching a paper airplane: you don't have any control over its trajectory or its landing spot.

would like to have more controls on our training processes

A **callback** is an object (a class instance implementing specific methods) that is passed to the model in the call to `fit()`

Inside of the fit(), can define callback
→ It can make control e.g) early stopping.

It has access to all the available data about the state of the model and its performance, and it can take action: interrupt training, save a model, load a different weight set, or otherwise alter the state of the model.



Using callbacks

Early stopping & Model checkpoint are widely used.

Model checkpointing—Saving the current state of the model at different points during training. `keras.callbacks.ModelCheckpoint`

Early stopping—Interrupting training when the validation loss is no longer improving (saving the best model obtained during training). `keras.callbacks.EarlyStopping`

Dynamically adjusting the value of certain parameters during training—Such as the learning rate of the optimizer. `keras.callbacks.LearningRateScheduler`, `keras.callbacks.ReduceLROnPlateau`

* callbacks related to visualizing

Logging training and validation metrics during training, or visualizing the representations learned by the model as they're updated. `keras.callbacks.CSVLogger`



The **EarlyStopping** callback interrupts training once a target metric being monitored has stopped improving for a fixed number of epochs

ModelCheckpoint lets you continually save the model during training

Listing 7.19 Using the callbacks argument in the fit() method

```
Callbacks are passed to the model via the
callbacks argument in fit(), which takes a list of
callbacks. You can pass any number of callbacks.
    ↗ we have to callbacks. EarlyStopping and Modelcheckpoint
callbacks_list = [
    keras.callbacks.EarlyStopping(
        monitor="val_accuracy",
        patience=2,
    ),
    keras.callbacks.ModelCheckpoint(
        filepath="checkpoint_path.keras",
        ↗ set file path to save the model
        monitor="val_loss",
        save_best_only=True,
    )
]
model = get_mnist_model()
model.compile(optimizer="rmsprop",
              loss="sparse_categorical_crossentropy",
              metrics=["accuracy"])
model.fit(train_images, train_labels,
          epochs=10,           ↗ I can use callback in fit stage.
          callbacks=callbacks_list,
          validation_data=(val_images, val_labels))
```

Saves the current weights after every epoch

Path to the destination model file

Interrupts training when improvement stops

Monitors the model's validation accuracy

Interrupts training when accuracy has stopped improving for two epochs

These two arguments mean you won't overwrite the model file unless val_loss has improved, which allows you to keep the best model seen during training.

You monitor accuracy, so it should be part of the model's metrics.

Note that because the callback will monitor validation loss and validation accuracy, you need to pass validation_data to the call to fit().



Writing your own callbacks

simple example that saves a list of per-batch loss values during training and saves a graph of these values at the end of each epoch.

Listing 7.20 Creating a custom callback by subclassing the `Callback` class

```
from matplotlib import pyplot as plt

class LossHistory(keras.callbacks.Callback):
    def on_train_begin(self, logs):
        self.per_batch_losses = [] # make list space

    def on_batch_end(self, batch, logs):
        self.per_batch_losses.append(logs.get("loss"))

    def on_epoch_end(self, epoch, logs):
        plt.clf()
        plt.plot(range(len(self.per_batch_losses)), self.per_batch_losses,
                 label="Training loss for each batch")
        plt.xlabel(f"Batch (epoch {epoch})")
        plt.ylabel("Loss")
        plt.legend()
        plt.savefig(f"plot_at_epoch_{epoch}")
        self.per_batch_losses = []
```

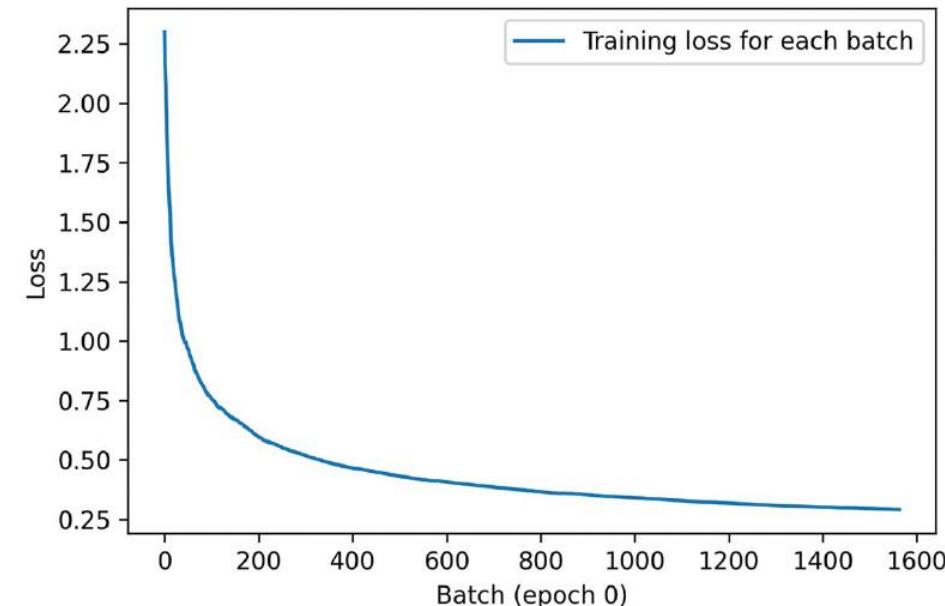
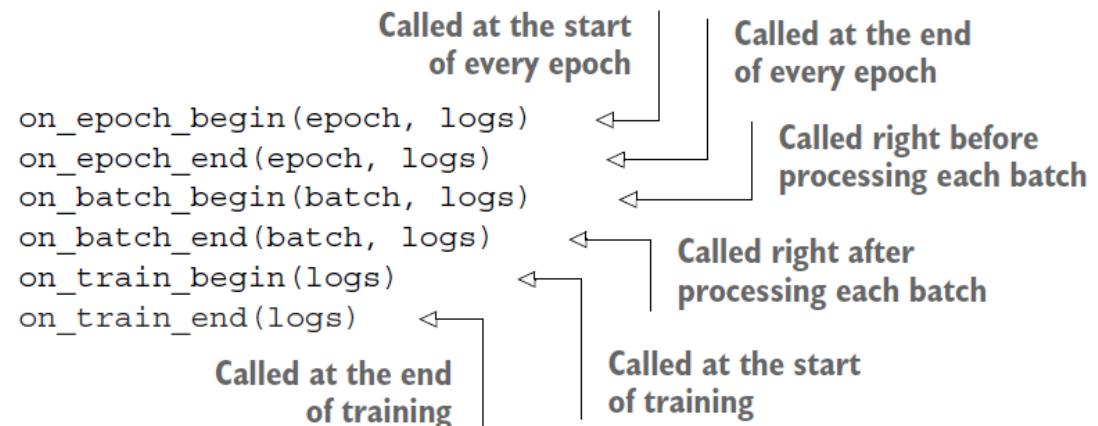


Figure 7.5 The output of our custom history plotting callback





chapter07_working-with-keras.i

File Edit View Insert Runtime Tools Help Cannot save changes

Share



Table of contents

- Subclassing the Model class
 - Rewriting our previous example as a subclassed model
 - Beware: What subclassed models don't support
- Mixing and matching different components
- Remember: Use the right tool for the job
- Using built-in training and evaluation loops
- Writing your own metrics
- Using callbacks
 - The EarlyStopping and ModelCheckpoint callbacks**
 - Writing your own callbacks
 - Monitoring and visualization with TensorBoard
 - Writing your own training and evaluation loops
 - Training versus inference
 - Low-level usage of metrics
 - A complete training and

+ Code + Text Copy to Drive

RAM Disk

Editing

▼ The EarlyStopping and ModelCheckpoint callbacks

Using the callbacks argument in the fit() method

```
callbacks_list = []
    keras.callbacks.EarlyStopping(
        monitor="val_accuracy",
        patience=2,
    ),
    keras.callbacks.ModelCheckpoint(
        filepath="checkpoint_path.keras",
        monitor="val_loss",
        save_best_only=True,
    )
]
model = get_mnist_model()
model.compile(optimizer="rmsprop",
              loss="sparse_categorical_crossentropy",
              metrics=["accuracy"])
model.fit(train_images, train_labels,
          epochs=10,
          callbacks=callbacks_list,
          validation_data=(val_images, val_labels))
```

```
[ ] model = keras.models.load_model("checkpoint_path.keras")
```

✓ 23s completed at 3:57 PM



검색하려면 여기에 입력하십시오.



16°C 맑음 오후 3:57
2022-04-03

Monitoring and visualization with TensorBoard

TensorBoard (www.tensorflow.org/tensorboard) is a browser-based application that you can run locally

Monitor training process of our model.

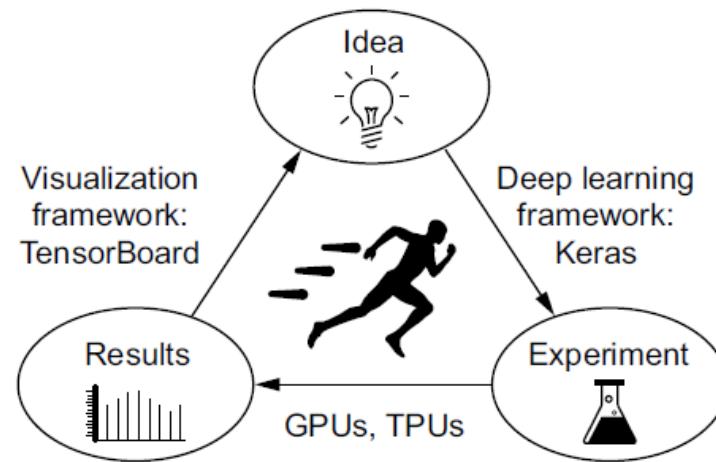


Figure 7.6 The loop of progress

- **Visually monitor metrics during training**
- **Visualize your model architecture**
- **Visualize histograms of activations and gradients**
- **Explore embeddings in 3D**



Easiest way to use TensorBoard with a Keras model and the fit() method is to use the `keras.callbacks.TensorBoard` callback.

```
model = get_mnist_model()
model.compile(optimizer="rmsprop",
              loss="sparse_categorical_crossentropy",
              metrics=["accuracy"])

tensorboard = keras.callbacks.TensorBoard(
    log_dir="/full_path_to_your_log_dir",
)
model.fit(train_images, train_labels,
          epochs=10,
          validation_data=(val_images, val_labels),
          callbacks=[tensorboard])
```

If you are running your script in a Colab notebook

```
%load_ext tensorboard
%tensorboard --logdir /full_path_to_your_log_dir
```

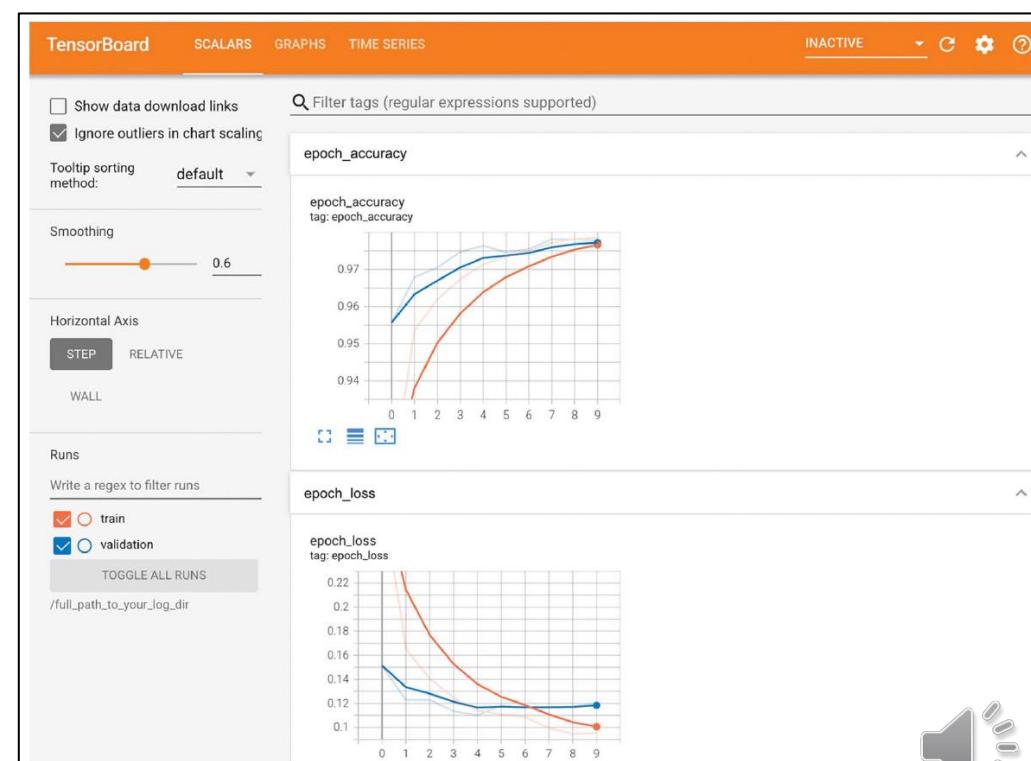


Figure 7.7 TensorBoard can be used for easy monitoring of training and evaluation metrics.



colab.research.google.com/github/fchollet/deep-learning-with-python-notebooks/blob/master/chapter07_working-with-keras.ipynb#scrollTo=5hNizEX2O8Sh

news study data science 학회사이트 업무관련 기타 Bookmark Manager Gmail YouTube 지도 뉴스 번역 SR Translate scm Data 학술지관련 » 기타

Share

chapter07_working-with-keras.i

File Edit View Insert Runtime Tools Help Cannot save changes

Table of contents

- Subclassing the Model class
- Rewriting our previous example as a subclassed model
- Beware: What subclassed models don't support
- Mixing and matching different components
- Remember: Use the right tool for the job
- Using built-in training and evaluation loops
- Writing your own metrics
- Using callbacks
- The EarlyStopping and ModelCheckpoint callbacks

Writing your own callbacks

- Monitoring and visualization with TensorBoard
- Writing your own training and evaluation loops
- Training versus inference
- Low-level usage of metrics
- A complete training and

+ Code + Text Copy to Drive

```
plot_at_epoch_0.png,
'plot_at_epoch_0.png',
'plot_at_epoch_2.png',
'sample_data']
```

RAM Disk

Monitoring and visualization with TensorBoard

```
[ ] model = get_mnist_model()
model.compile(optimizer="rmsprop",
              loss="sparse_categorical_crossentropy",
              metrics=["accuracy"])

tensorboard = keras.callbacks.TensorBoard(
    log_dir="/full_path_to_your_log_dir",
)
model.fit(train_images, train_labels,
          epochs=10,
          validation_data=(val_images, val_labels),
          callbacks=[tensorboard])
```

```
[ ] %load_ext tensorboard
%tensorboard --logdir /full_path_to_your_log_dir
```

Writing your own training and evaluation loops

Training versus inference

0s completed at 4:03 PM

검색하려면 여기에 입력하십시오.

15°C 맑음

오늘 4:08
2022-04-03

Writing your own training and evaluation loops

Built-in fit() workflow is solely focused on supervised learning, However, not every form of machine learning falls into this category.

Typical training loop look like this:

1. Run the forward pass inside a gradient tape to obtain a loss for the current batch of data.
2. Retrieve the gradients of the loss with regard to the model's weights.
3. Update the model's weights so as to lower the loss value on the current batch of data.

We will learn to reimplement fit()



Supervised-learning training step ends up looking like this:

```
def train_step(inputs, targets):  
    with tf.GradientTape() as tape:  
        predictions = model(inputs, training=True) *  
        loss = loss_fn(targets, predictions)  
    gradients = tape.gradients(loss, model.trainable_weights)  
    optimizer.apply_gradients(zip(model.trainable_weights, gradients))
```



Set training = True: some Keras layers, such as the Dropout layer, have different behaviors during training and during inference

Layers and models own two kinds of weights: trainable weights, non-trainable weights

→ weights that automatically updated
No need to update myself

↑
only need to update
trainable weights.



A complete training and evaluation loop

Training step function takes a batch of data and targets

Listing 7.21 Writing a step-by-step training loop: the training step function

```

model = get_mnist_model()

loss_fn = keras.losses.SparseCategoricalCrossentropy()
optimizer = keras.optimizers.RMSprop()
metrics = [keras.metrics.SparseCategoricalAccuracy()]
loss_tracking_metric = keras.metrics.Mean()

def train_step(inputs, targets):
    with tf.GradientTape() as tape:
        predictions = model(inputs, training=True)
        loss = loss_fn(targets, predictions)
        gradients = tape.gradient(loss, model.trainable_weights)
        optimizer.apply_gradients(zip(gradients, model.trainable_weights))

    logs = {}
    for metric in metrics:
        metric.update_state(targets, predictions)
        logs[metric.name] = metric.result()

    loss_tracking_metric.update_state(loss)
    logs["loss"] = loss_tracking_metric.result()

    return logs
  
```

How we defining training step

Prepare the loss function.

Prepare the optimizer.

Prepare the list of metrics to monitor.

Prepare a Mean metric tracker to keep track of the loss average.

Run the forward pass. Note that we pass training=True.

Run the backward pass. Note that we use model.trainable_weights.

Keep track of metrics.

Keep track of the loss average.

Return the current values of the metrics and the loss.

we get the mean value of (batch samples) losses.

eg) 1000 data point
batch size 10
100 batches
epoch 0 → training updates 100 times
for each batch, it contains (number of samples)
train_step will apply for each batch



Reset the state of our metrics at the start of each epoch and before running evaluation.

Listing 7.22 Writing a step-by-step training loop: resetting the metrics

```
def reset_metrics():
    for metric in metrics:
        metric.reset_state()
loss_tracking_metric.reset_state()
```

Our complete training loop

Listing 7.23 Writing a step-by-step training loop: the loop itself

```
training_dataset = tf.data.Dataset.from_tensor_slices(
    (train_images, train_labels))
training_dataset = training_dataset.batch(32)
epochs = 3
for epoch in range(epochs):
    reset_metrics()
    for inputs_batch, targets_batch in training_dataset:
        logs = train_step(inputs_batch, targets_batch)
    print(f"Results at the end of epoch {epoch}")
    for key, value in logs.items():
        print(f"...{key}: {value:.4f}")
```



The evaluation loop

Listing 7.24 Writing a step-by-step evaluation loop

```
def test_step(inputs, targets): important
    predictions = model(inputs, training=False)
    loss = loss_fn(targets, predictions)

    logs = {} evaluation metrics & loss
    for metric in metrics:
        metric.update_state(targets, predictions)
        logs["val_" + metric.name] = metric.result()
    loss_tracking_metric.update_state(loss)
    logs["val_loss"] = loss_tracking_metric.result()
    return logs

val_dataset = tf.data.Dataset.from_tensor_slices((val_images, val_labels))
val_dataset = val_dataset.batch(32)
reset_metrics()
for inputs_batch, targets_batch in val_dataset:
    logs = test_step(inputs_batch, targets_batch)
print("Evaluation results:")
for key, value in logs.items():
    print(f"...{key}: {value:.4f}")
```

Note that we pass
training=False.



Make it fast with `tf.function`

Listing 7.25 Adding a `@tf.function` decorator to our evaluation-step function

```
    ↗ using GPU, running codes below
@tf.function
def test_step(inputs, targets):
    predictions = model(inputs, training=False)
    loss = loss_fn(targets, predictions)

    logs = {}
    for metric in metrics:
        metric.update_state(targets, predictions)
        logs["val_" + metric.name] = metric.result()

    loss_tracking_metric.update_state(loss)
    logs["val_loss"] = loss_tracking_metric.result()
    return logs

val_dataset = tf.data.Dataset.from_tensor_slices((val_images, val_labels))
val_dataset = val_dataset.batch(32)
reset_metrics()

for inputs_batch, targets_batch in val_dataset:
    logs = test_step(inputs_batch, targets_batch)
print("Evaluation results:")
for key, value in logs.items():
    print(f"...{key}: {value:.4f}")
```

This is the
only line that
changed.

On the Colab CPU, we go from taking 1.80 s to run the evaluation loop to only 0.8 s. 

Leveraging fit() with a custom training loop

complicated

There's actually a middle ground between fit() and a training loop written from scratch: you can provide a custom training step function and let the framework do the rest.

Listing 7.26 Implementing a custom training step to use with fit()

```
loss_fn = keras.losses.SparseCategoricalCrossentropy()
loss_tracker = keras.metrics.Mean(name="loss")           ← This metric object will be used to
                                                        track the average of per-batch losses
                                                        during training and evaluation.

class CustomModel(keras.Model):
    def train_step(self, data):      ← We override the
                                    train_step method.
        inputs, targets = data
        with tf.GradientTape() as tape:
            predictions = self(inputs, training=True)
            loss = loss_fn(targets, predictions)
        gradients = tape.gradient(loss, model.trainable_weights)
        optimizer.apply_gradients(zip(gradients, model.trainable_weights))

        loss_tracker.update_state(loss)
        return {"loss": loss_tracker.result()}           ← We update the loss
                                                        tracker metric that
                                                        tracks the average
                                                        of the loss.

@property
def metrics(self):
    return [loss_tracker]                         ← Any metric you
                                                would like to reset
                                                across epochs should
                                                be listed here.

                                                ← We return the average loss
                                                so far by querying the loss
                                                tracker metric.
```



colab.research.google.com/github/fchollet/deep-learning-with-python-notebooks/blob/master/chapter07_working-with-keras.ipynb#scrollTo=aoT8OjHm6ybL

news study data science 학회사이트 업무관련 기타 Bookmark Manager Gmail YouTube 지도 뉴스 번역 SR Translate scm Data 학술지관련 » 기타

Share

Table of contents

- Remember: Use the right tool for the job
- Using built-in training and evaluation loops
 - Writing your own metrics
 - Using callbacks
 - The EarlyStopping and ModelCheckpoint callbacks
 - Writing your own callbacks
 - Monitoring and visualization with TensorBoard
- Writing your own training and evaluation loops
 - Training versus inference
- Low-level usage of metrics**
 - A complete training and evaluation loop
 - Make it fast with tf.function
 - Leveraging fit() with a custom training loop
- Summary

+ Code + Text Copy to Drive RAM Disk Editing

Training versus inference

Low-level usage of metrics

```
metric = keras.metrics.SparseCategoricalAccuracy()  
targets = [0, 1, 2]  
predictions = [[1, 0, 0], [0, 1, 0], [0, 0, 1]]  
metric.update_state(targets, predictions)  
current_result = metric.result()  
print(f"result: {current_result:.2f}")
```

result: 1.00

```
[16] values = [0, 1, 2, 3, 4]  
mean_tracker = keras.metrics.Mean()  
for value in values:  
    mean_tracker.update_state(value)  
print(f"Mean of values: {mean_tracker.result():.2f}")
```

Mean of values: 2.00

A complete training and evaluation loop

Writing a step-by-step training loop: the training step function

0s completed at 4:56 PM

검색하려면 여기에 입력하십시오.

15°C 맑음

오늘 4:56
2022-04-03

Thank you!